

T3: Introduction to Data Analytics

Data Analytics with HPC
Master in High Performance Computing, Spring 2026



Contents

- 12/2 - Introduction to Python
 - Execution environments:
iPython, Jupyter.
 - Basics.
 - Operators and types.
Instructions. Functions.
Modules. Classes. Exceptions.
 - Integrated tools.
 - NumPy.
 - Matplotlib.

Contents

- 19/2 - Data Analytics: Pandas
 - Intro to Pandas.
 - Data I/O.
 - Data wrangling.
 - Visualization.
 - Aggregation.
 - Time series.

Contents

- 26/2 - High level tools
 - Visualization: Seaborn.
 - Machine learning: Scikit-Learn
 - Out-of-core computing: Dask

Contents

- 5/3 - Projects (**online only**)
 - Time for developing your data analysis project.

Introduction to Python

Contents


- Execution environments: iPython, Jupyter.
- Basics.
- Operators and types. Instructions. Functions. Modules. Classes. Exceptions.
- Integrated tools.
- NumPy.
- Matplotlib.

- General-purpose language, interpreted, object-oriented.
- Portable: an interpreter exists for almost any OS.
- Open-source reference implementation (CPython), but alternative ones exist as well (e.g., iPython, Intel Python).
- Powerful and versatile:
 - Dynamic typing.
 - Complex “basic” types: lists, dictionaries, strings, ...
 - Built-in tools: powerful list operations, strings, arrays, ...
 - Libraries: math, statistics, parallelization, ...
 - Automatic memory management.

- Popular with the scientific and engineering communities.
- Free and general purpose alternative to Matlab/R.
- “Slice” notation (Matlab-like).
- Multiple extensions: NumPy, Matplotlib, SciPy, pytables, ...
- Simple integration: C/C++/Fortran, web services, ...
- Web development: Django, Flask, Pyramid, ...
- Scientific computing environments: iPython, conda-forge, ...


Python: “state of the art”

Which version of Python do you use the most?



One in four respondents use the latest version of Python, released in October 2022. It took Python 3.11 approximately nine months to reach 27% adoption.

What do you use Python for?



Show more

The use cases of Python have remained stable year over year, with the most popular areas of usage being data analysis, web development, and machine learning.

Source: <https://www.jetbrains.com/lp/devecosystem-2023/python/>


- Interactive Python interpreter.
- Does not provide analytics or computing tools, but greatly eases up the use of Python and the development process.
- Includes GUIs for easy data visualization.
- Allows to work remotely through a web browser, as well as storing work sessions in an HTML format to allow sharing code, data, and examples.

Jupyter Notebooks



- Environment focused on exporting HTML or being used interactively.
- Uses a JSON-based format to allow sharing code and outputs.
- Built as a web server to which clients are connected using a browser.
- Allows to remotely use a Python environment via web.

Jupyter Notebooks





Jupyter Notebooks

The screenshot shows a Jupyter Notebook interface running in a Chromium browser. The title bar reads "Untitled - Chromium". The address bar shows the URL "localhost:8888/notebooks/Untitled.ipynb?kernel_name=python2". The browser's toolbar includes standard icons for back, forward, search, and refresh.

The Jupyter interface has a header with the title "jupyter Untitled Last Checkpoint: a minute ago (unsaved changes)" and a Python 2 kernel icon. Below the header is a toolbar with various icons for file operations like new, open, save, and cell execution.

The main area contains three code cells:

```
In [1]: %pylab inline
Populating the interactive namespace from numpy and matplotlib
```

```
In [2]: img=plt.imread( '/tmp/Mandelbrot-Javascript.png' )
```

```
In [3]: imshow(img)
Out[3]: <matplotlib.image.AxesImage at 0x7f847692b410>
```

Cell In [3] displays a plot of the Mandelbrot set, rendered in red and black against a white background. The plot is a fractal with a large central black region and smaller red lobes extending from it. The axes range from 0 to 600 on both the x and y axes.

Google Colaboratory



- Free notebook environment on the cloud.
- <https://colab.research.google.com>

The screenshot shows the Google Colaboratory interface. At the top, there's a navigation bar with a back arrow, forward arrow, refresh button, and a URL bar containing "colab.research.google.com/notebooks/welcome.ipynb". Below the bar is the main menu with options like File, Edit, View, Insert, Runtime, Tools, Help, and a toolbar with various icons. On the left, a sidebar displays a "Table of contents" with sections such as "Introducing Colaboratory", "Getting Started", "More Resources", and "Machine Learning Examples: Seedbank". A "Code snippets" and "Files" tab are also visible. The main content area features a heading "Welcome to Colaboratory!" followed by a brief introduction: "Colaboratory is a free Jupyter notebook environment that requires no setup and runs entirely in the cloud. With Colaboratory you can write and execute code, save and share your analyses, and access powerful computing resources, all for free from your browser." Below this, a section titled "Introducing Colaboratory" contains a video thumbnail for "Get started with Google Colaboratory (Coding)". The video thumbnail shows a man smiling and includes the text "Intro to Google Colab" and "Coding TensorFlow". At the bottom of the content area, there's a section titled "Getting Started" with a note about the document being a Jupyter notebook.

Welcome To Colaboratory - C X +

colab.research.google.com/notebooks/welcome.ipynb

CO Welcome To Colaboratory

File Edit View Insert Runtime Tools Help

+ Code + Text ⌘ Copy to Drive

Share G

Connect Editing

Table of contents Code snippets Files

Introducing Colaboratory

Getting Started

More Resources

Machine Learning Examples: Seedbank

Section

Welcome to Colaboratory!

Colaboratory is a free Jupyter notebook environment that requires no setup and runs entirely in the cloud. With Colaboratory you can write and execute code, save and share your analyses, and access powerful computing resources, all for free from your browser.

Introducing Colaboratory

This 3-minute video gives an overview of the key features of Colaboratory:

Get started with Google Colaboratory (Coding)

Intro to Google Colab

Coding TensorFlow

Getting Started


The document you are reading is a [Jupyter notebook](#), hosted in Colaboratory. It is not a static page, but an interactive environment that lets you write and execute code in Python and other languages.

Jupyter QtConsole


File Edit View Kernel Window Help
 kernel 0 (kernel 0) slave

```
Jupyter QtConsole 4.6.0
Python 3.8.1 | packaged by conda-forge | (default, Jan 5 2020, 20:58:18)
Type 'copyright', 'credits' or 'license' for more information
IPython 7.11.1 -- An enhanced Interactive Python. Type '?' for help.

In [14]: plt.scatter(X[:,0],X[:,1], c=y, s=50, cmap='Spectral')
Out[14]: <matplotlib.collections.PathCollection at 0x7f214f5dde50>
```




```
In [15]: fig, ax = plt.subplots(1, 4, figsize=(16, 3))
...: fig.subplots_adjust(left=0.02, right=0.98, wspace=0.1)
...: for axi, depth in zip(ax, range(1, 5)):
...:     model = DecisionTreeClassifier(max_depth=depth)
...:     visualize_tree(model, X, y, ax=axi)
...:     axi.set_title('depth = {}'.format(depth))
...:
<ipython-input-7-ac8555acc73>:24: UserWarning: The following kwargs were not used by contour: 'clim'
contours = ax.contourf(xx, yy, Z, alpha=0.3,
<ipython-input-7-ac8555acc73>:24: UserWarning: The following kwargs were not used by contour: 'clim'
contours = ax.contourf(xx, yy, Z, alpha=0.3,
<ipython-input-7-ac8555acc73>:24: UserWarning: The following kwargs were not used by contour: 'clim'
contours = ax.contourf(xx, yy, Z, alpha=0.3,
<ipython-input-7-ac8555acc73>:24: UserWarning: The following kwargs were not used by contour: 'clim'
contours = ax.contourf(xx, yy, Z, alpha=0.3,
```



```
In [16]: |
```

JupyterLab: Next Generation Notebook Interface 



Python: Conceptual hierarchy

- Programs are made of **modules**.
- Modules contain **functions**.
- Functions contain **instructions**.
- Instructions contain **expressions**.
- Expressions create and manipulate **objects**.
- Objects belong to a **class**.

Python: Conceptual hierarchy

```
>>>>>> file.py
```

```
def f():
    i = 0
```

```
while i < 100:
    print( i )
    i += 1
```

Python: Objects

- **Everything** is an object:
 - integer numbers -> *int*
 - floating point numbers -> *float*
 - character strings -> *str*
 - lists -> *list*
 - dictionaries -> *dict*
- The type of a variable can be queried using *type()*.
- Objects are containers which aggregate variables (*attributes*) and functions (*methods*).
- Objects belong to a *class*, which is an abstract description (or scheme) of the common traits of a family of objects.

Python - Executing programs



1. Using the interactive interpreter:

```
>>> txt = "Luke, I am your father"  
>>> print( txt )  
Luke, I am your father
```

Python - Executing programs



2. Interpreting Python code:

```
----- > script.py
```

```
txt = "Luke, I am your father"  
print( txt )
```

```
-----
```

```
$ python script.py
```

```
Luke, I am your father
```

Python - Executing programs



3. Shell script:

```
----- > script.py
```

```
#! /usr/bin/python

txt = "Luke, I am your father"

print( txt )
```

```
-----
```

```
$ ./script.py
```

```
Luke, I am your father
```

Python - Executing programs



4. Embedding Python into C:

```
----- > test.c

#include <Python.h>

...
Py_Initialize();

PyRun_SimpleString( "txt = 'Luke, I am your father'" );
PyRun_SimpleString( "print( txt )" );
```

...

```
-----
```



```
$ gcc -I /usr/include/python3.8 test.c -lpython3.8
$ ./a.out
```

Luke, I am your father

Python - Semantics & Syntax

Indentation, not brackets



```
if x < threshold:  
    print( "x less than threshold" )  
  
else:  
    print( "x greater than threshold" )
```

Python - Semantics & Syntax

Everything is an object



- Including numbers, strings, functions, classes, modules, ...
- Each object has a type and its own data.
- This provides flexibility, allowing to deal with any element in a generic way.

Python - Semantics & Syntax

Comments



```
# This is a dramatic moment.  
  
txt = "Luke, I am your father"  
  
print( txt )
```

Python - Semantics & Syntax

Function calls



- C syntax:

```
result = f( x, y, z )
```

- Class methods:

```
result = obj.f( x, y, z )
```

- Functions may be passed positional or keyword arguments:

```
result = f( x, y, z, tol=0.01, method="fast" )
```

Python - Semantics & Syntax

Variables and passing by reference

- Assignment to a variable generates a new reference to an object.
- The original object remains the same!

```
>>> a = [1, 2, 3]
```

```
>>> b = a
```

```
>>> a.append(4)
```

```
>>> b
```

```
[1, 2, 3, 4]
```

Python - Semantics & Syntax

Variables and passing by reference

Consequently:

- Passing large variables is efficient.
- Functions may permanently modify their parameters!
- Except for **immutable** types. E.g.,

```
>>> a = 10
```

```
>>> b = a
```

```
>>> a = 20
```

```
>>> print b
```

10

Python - Semantics & Syntax

Dynamic references, strong typing



References (variables) have no static type:

```
>>> a = 10  
  
>>> type(a)  
  
<type 'int'>  
  
>>> a = "10"  
  
>>> type(a)  
  
<type 'str'>
```

Python - Semantics & Syntax

Dynamic references, strong typing



Python has strong typing, nevertheless:

```
>>> 5 + "5"
```

```
TypeError: unsupported operand type(s) for +: 'int' and 'str'
```

```
>>> "5" + 5
```

```
TypeError: cannot concatenate 'str' and 'int' objects
```

```
>>> 5 + int("5")
```

```
10
```

```
>>> "5" + str(5)
```

```
'55'
```

Each Python object has a specific type (class). Implicit casting is only performed for `int` to floating point conversions.

Python - Semantics & Syntax

Dynamic references, strong typing

The `isinstance` operator is used to check whether an object belongs to a particular class:

```
>>> a = 4.5
```

```
>>> isinstance(a, int)
```

```
False
```

```
>>> isinstance(a, float)
```

```
True
```



Python - Semantics & Syntax

Methods and Attributes

Python objects have methods and attributes:

```
>>> a = 4.5
```

```
>>> dir(a)
```

```
['__abs__', '__add__', '__class__', '__coerce__', '__delattr__', '__div__', '__divmod__',
 '__doc__', '__eq__', '__float__', '__floordiv__', '__format__', '__ge__',
 '__getattribute__', '__getformat__', '__getnewargs__', '__gt__', '__hash__',
 '__init__', '__int__', '__le__', '__long__', '__lt__', '__mod__', '__mul__',
 '__ne__', '__neg__', '__new__', '__nonzero__', '__pos__', '__pow__',
 '__radd__', '__rdiv__', '__rdivmod__', '__reduce__', '__reduce_ex__',
 '__repr__', '__rfloordiv__', '__rmod__', '__rmul__',
 '__rpow__', '__rsub__', '__rtruediv__', '__setattr__', '__setformat__',
 '__sizeof__', '__str__', '__sub__', '__subclasshook__', '__truediv__',
 '__trunc__', 'as_integer_ratio',
 'conjugate', 'fromhex', 'hex', 'imag', 'is_integer', 'real']
```

Python - Semantics & Syntax

Methods and Attributes

```
>>> a.hex()  
  
'0x1.200000000000p+2'  
  
>>> hasattr( a, "hex" )  
  
True  
  
>>> getattr( a, "hex" )  
  
<built-in method hex of float object at 0x1ed86a0>  
  
>>> getattr( a, "hex" )()  
  
'0x1.200000000000p+2'
```

Python - Semantics & Syntax

Importing modules



- A Python module is a .py file including definitions and/or code:

```
----- > modulo1.py
```

```
PI = 3.14159
```

```
def f(x):  
    return x + 2
```

```
def g(a,b):  
    return a + b
```

```
-----
```

Python - Semantics & Syntax

Binary operators



<code>a + b</code>	Addition	<code>a - b</code>	Subtraction
<code>a * b</code>	Multiplication	<code>a / b</code>	Division (Python 3)
<code>a // b</code>	Integer division (Python 3)	<code>a ** b</code>	Exponentiation
<code>a & b</code>	Bitwise AND	<code>a b</code>	Bitwise OR
<code>a ^ b</code>	Bitwise XOR	<code>a == b</code>	Equality
<code>a != b</code>	Inequality	<code>a < b, a <= b</code>	Less/Less or equal than
<code>a > b, a >= b</code>	Greater/Greater or equal than	<code>a is b</code>	Identity: True if <code>a</code> and <code>b</code> are the same object

Python - Scalar types



None	“null” in Python. Singleton class.
str	String type. ASCII in Python 2.x, Unicode in Python 3.
unicode	Unicode string.
float	Floating point number, double precision (64 bits).
bool	Logical, True or False.
int	Signed integer. 32 or 64 bits depending on platform. Arbitrary precision in Python 3.
long	Signed integer, arbitrary precision. Does not exist in Python 3.

Python - Scalar types

Numerical types

- Python 2: `int` (32- or 64-bits), `float` (64 bits) and `long` (arbitrary precision).
- Python 3: `int` (arbitrary precision), `float` (64 bits).
- Conversion between `int` and `long` is transparent in Python 2.
- Scientific notation accepted:

```
>>> fval = 6.78e-5
```

Python - Scalar types

Numerical types

- Division of two integers in Python 2 yields an integer (truncated, not rounded). For floating point division of integers in Python 2:

```
>>> fval = a_int / float(b_int)
```

- Division of integers in Python 3 yields a floating point number. For integer division of integers in Python 3:

```
>>> ival = a_int // b_int
```

- Complex numbers are written using `j` for the imaginary part:

```
>>> cval = 1 + 2j
```

Python - Scalar types

Strings

- String literals are written between single or double quotes:

```
>>> a = 'one way to write a string'
```

```
>>> b = "a different way"
```

- To write strings with line breaks, use three single quotes:

```
>>> c = '''
```

This is a larger string

Which spans more than one line

```
'''
```

Python - Scalar types

Strings

- Strings are immutable objects: once they are created, they cannot be modified:

```
>>> a = 'this is a string'
```

```
>>> a[10] = 'f'
```

```
TypeError: 'str' object does not support item assignment
```

```
>>> b = a.replace('string', 'larger string')
```

```
>>> b
```

```
'this is a larger string'
```

Python - Scalar types

Strings

- Strings can be created from different types using `str()`:

```
>>> a = 5.6
```

```
>>> s = str(a)
```

```
>>> s
```

```
'5.6'
```

Python - Scalar types

Strings

- A string is just a **collection** of characters (so not really a scalar type):

```
>>> s = 'python'  
  
>>> list(s)  
  
[ 'p', 'y', 't', 'h', 'o', 'n' ]  
  
>>> s[:3]  
  
'pyt'
```

Python - Scalar types

Strings

- The backslash (\) is used as the escape character, to encode special characters such as line breaks (\n) or Unicode characters.
- A string which should be interpreted literally (raw) can be marked using r'string':

```
>>> s = r'no\special\characters\in>this\string'
```

```
>>> s
```

```
'no\special\characters\in>this\string'
```

Python - Scalar types

Strings

- Operator + applies string concatenation:

```
>>> a = 'first half of string '
```

```
>>> b = 'and second half of string'
```

```
>>> a + b
```

```
'First half of string and second half of string'
```

Python - Scalar types

Strings

- To specify that a string is encoded using Unicode, `u' string'` is used.
- To format a string there's a C-like template syntax:

```
>>> template = u"% .2f %s are worth %d€"
```

```
>>> template % (1.18018, u'americian dollars', 1)
```

```
u'1.18 american dollars are worth 1\u20ac'
```

```
>>> print template % (1.18018, u'americian dollars', 1)
```

```
'1.18 american dollars are worth 1€'
```

Python - Scalar types

Booleans

- The two boolean values are written `True` and `False`.
- They can be combined using the keywords `and` and `or`.

```
>>> True and True
```

```
True
```

```
>>> False and True
```

```
True
```

Python - Scalar types

Booleans

- All basic predefined types in Python, as well as any class implementing the `__nonzero__()` method, can be interpreted as `True` or `False` in a conditional clause.

```
>>> a = [1,2,3]

>>> if a: print "True"
'True'

>>> b = []

>>> if not b: print "False"
'False'
```

Python - Scalar types

None

- `None` is the null value in Python.
- If a function does not explicitly return a value, it returns `None` implicitly.
- Can be used as a default value for optional parameters to functions:

```
def add_and_maybe_multiply( a, b, c = None ):  
  
    result = a + b  
  
    if c is not None: result *= c  
  
    return result
```

- `None` is not a reserved word, but an instance (singleton) of `NoneType`.

Python - Scalar types

Dates and times

- The `datetime` provides the types `datetime`, `date` and `time`.

```
>>> from datetime import datetime, date, time
```

```
dt = datetime( 2015, 04, 24, 12, 25, 32 )
```

```
>>> dt.day
```

```
24
```

```
>>> dt.time()
```

```
datetime.time( 12, 25,  
32 )
```

Python - Scalar types

Dates and times

- The `strftime()` method formats a `datetime` object as a string:

```
>>> dt.strftime( '%d/%m/%Y %H:%M' )
```

```
'24/04/2015 12:25'
```

- A `datetime` object can be created from a string using `strptime`:

```
>>> datetime.strptime( '20150424', '%Y%m%d' )
```

```
datetime.datetime( 2015, 4, 24, 0, 0 )
```

Python - Scalar types

Dates and times

- The subtraction of two `datetime` objects returns a `datetime.timedelta`:

```
>>> dt2 = datetime( 2015, 4, 25, 12, 25, 32 )
```

```
>>> dt2 - dt
```

```
datetime.timedelta(1)
```

- Operating a `datetime` object with a `timedelta` returns a new `datetime`:

```
>>> from datetime import timedelta
```

```
>>> dt + timedelta(days=1)
```

```
>>> datetime.datetime(2015, 4, 25, 12, 25, 32)
```

Python - Flow control

if, elif, else

These blocks behave like in most other languages:

```
if x < 0:  
    print 'Negative'  
  
elif x == 0:  
    print 'Zero'  
  
elif 0 < x < 5:  
    print 'Positive less than 5'  
  
else:  
    print 'Positive and greater or equal than 5'
```

Python - Flow control

for loops

- A for loop is used to iterate over a collection (like a list or tuple):

```
for val in collection:
```

```
    # Do something with val
```

- continue skips to the next iteration
- break exits the current loop

Python - Flow control

range() and xrange()

- range() returns a list of equally spaced integers:

```
>>> range(10)
```

```
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
```

- Takes as parameters the start, end, and step of the sequence:

```
>>> range(5,20,3)
```

```
[5,8,11,14,17]
```

- Generates values in the interval [start, end).

Python - Flow control

range() and xrange()

- range() is useful to code for loops with C semantics:

```
>>> for x in range(10):
```

```
...     print 2*x
```

```
0
```

```
2
```

```
4
```

```
8
```

```
16
```

```
18
```

Python - Flow control

range() and xrange()

- xrange() accepts the same parameters as range(), but returns a generator instead of building the full list in memory.
- It features lazy evaluation.

```
>>> xrange(10)
```

```
xrange(10)
```

```
>>> for x in xrange(10):  
...     print 2*x
```

```
0
```

```
2
```

```
...
```

- It is preferable to use xrange() when working with large ranges.

Python - Flow control

while loops

- Iterate **while** a condition is met:

```
x = 256  
  
total = 0  
  
while x > 0:  
  
    if total > 500: break  
  
    total += x  
  
    x = x // 2
```

Python - Flow control

pass

- `pass` is a no-op instruction in Python.
- It is sometimes required because white spaces in Python delimit execution blocks:

```
if x < 0:  
    print "Negativo"  
  
elif x == 0: pass  
  
else:  
    print "Positivo"
```

Python - Flow control

Exception handling

- Functions and operations can raise exceptions:

```
>>> 5 / 0
```

```
ZeroDivisionError: integer division or modulo by zero
```

- Exceptions can be handled to solve runtime errors.
- This allows to dynamically manage some selected error types.

Python - Flow control

Exception handling

```
def float_division( a, b ):  
    x = NaN  
  
    try:  
  
        x = a / float(b)  
  
    except:  
        print "Exception during division"  
  
    return x
```

Python - Flow control

Exception handling

- We can explicitly list the exception types managed by a `except` block:

```
def float_division( a, b ):  
    x = NaN  
  
    try:  
  
        x = a / float(b)  
  
    except ZeroDivisionError:  
        print "Division by zero"  
  
    return x
```

Python - Flow control

Exception handling

- Multiple exception types can be managed by the same block:

```
def float_division( a, b ):  
  
    x = NaN  
  
    try:  
  
        x = a / float(b)  
  
    except ZeroDivisionError, TypeError:  
  
        print "Exception during division"  
  
    return x
```

Python - Flow control

Exception handling

- A finally block is executed regardless of whether the try was successful:

```
f = open( path, 'w' )

try:

    write_to_file( f )

finally:

    f.close()
```

Python - Flow control

Exception handling

- An `else` block will only be executed if the `try` was successful:

```
f = open( path, 'w' )

try:

    write_to_file( f )

except: print 'Failure'

else: print 'Success'

finally: f.close()
```

Python - Collections

Tuples



- Tuple: unidimensional sequence of fixed size containing Python objects.

```
>>> tup = (2, 3, 4)

>>> nested_tup = ( (1, 2), (2, 3, 4) )

>>> tup_from_list = tuple( [2,3,4] )

>>> tup_from_iter = tuple( xrange(5) )

>>> tup_from_string = tuple( 'string' )
```

Python - Collections

Tuples

- The elements in a tuple are accessed using the [] operator:

```
>>> tup_from_iter[3]
```

```
3
```

```
>>> tup_from_iter[1:3]
```

```
(1, 2)
```

```
>>> tup_from_string[0]
```

```
's'
```

Python - Collections

Tuples

- Tuples are immutable:

```
>>> tup_from_iter[3] = 5
```

```
TypeError: 'tuple' object does not support item assignment
```

- However, objects inside a tuple can be mutable:

```
>>> tup = ( (1,2), [3,4,5] )
```

```
>>> tup[1].append(6)
```

```
>>> tup
```

```
((1, 2), [3, 4, 5, 6])
```

Python - Collections

Tuples

- Tuples are concatenated using the + operator:

```
>>> tup_from_iter + tup_from_string  
  
(0, 1, 2, 3, 4, 's', 't', 'r', 'i', 'n', 'g')
```

- The * operator is consistent with the addition/concatenation semantics:

```
>>> tup_from_iter*3  
  
(0, 1, 2, 3, 4, 0, 1, 2, 3, 4, 0, 1, 2, 3, 4)
```

Python - Collections

Tuples

- Note that the objects contained in tuples are not copied, but referenced:

```
>>> tup = ( (1, 2), (3, 4, 5) )  
  
>>> tup2 = tup * 3  
  
>>> tup[1].append(6)  
  
>>> tup2  
((1, 2), [3, 4, 5, 6], (1, 2), [3, 4, 5, 6], (1, 2), [3, 4,  
5, 6])
```

Python - Collections

Tuples

- If the right hand side of an assignment is a tuple, Python tries to unpack it:

```
>>> tup = (1, 2, 3)
```

```
>>> a, b, c = tup
```

```
>>> b
```

Python - Collections

Tuples

- Nested tuples can be explicitly unpacked:

```
>>> tup = ( 1, 2, (3, 4) )
```

```
>>> a, b, (c, d) = tup
```

```
>>> c
```

```
3
```

```
>>> a, b = b, a
```

```
>>> a
```

```
2
```

- A common use of tuples is to code functions that return multiple values.

Python - Collections

Lists

- Unlike tuples, lists have variable length and are mutable:

```
>>> list_1 = [2, 3, 7, None]
```

```
>>> tup = ('a', 'b', 'c' )
```

```
>>> list_2 = list( tup )
```

```
>>> list_2[1] = 'd'
```

```
>>> list_2
```

```
[ 'a', 'd', 'c' ]
```

Python - Collections

Lists

- Elements are added at the end of a list using `append()`:

```
>>> list_2.append( 5 )
```

```
>>> list_2
```

```
[ 'a', 'd', 'c', 5 ]
```

- Elements can also be added at a specific place using `insert()`:

```
>>> list_2.insert( 3, None )
```

```
>>> list_2
```

```
[ 'a', 'd', 'c', None, 5 ]
```

Python - Collections

Lists

- The reverse operation to `insert()` is `pop()`:

```
>>> list_2.pop()
```

```
5
```

```
>>> list_2
```

```
[ 'a', 'd', 'c', None ]
```

```
>>> list_2.pop(2)
```

```
'c'
```

```
>>> list_2
```

```
[ 'a', 'd', None ]
```

Python - Collections

Lists

- Elements can be deleted from the list using `remove()`:

```
>>> list_2.append('a')
```

```
>>> list_2
```

```
[ 'a', 'd', None, 'a' ]
```

```
>>> list_2.remove('a')
```

```
>>> list_2
```

```
[ 'd', None, 'a' ]
```

Python - Collections

Lists

- The `in` operator checks whether a value is contained in a list:

```
>>> 'a' in list_2
```

True

```
>>> 'c' in list_2
```

False

Python - Collections

Lists

- The operator + is used to concatenate lists:

```
>>> list_1 + list_2  
[ 2, 3, 7, None, 'd', None, 'a' ]
```

- extend() adds full lists to a given one:

```
>>> list_1.extend( list_2 )  
  
>>> list_1  
[ 2, 3, 7, None, 'd', None, 'a' ]
```

Python - Collections

Lists

- A list can be sorted in place using `sort()`:

```
>>> a = [ 7, 2, 5, 1, 3 ]  
  
>>> a.sort()  
  
>>> a  
  
[1, 2, 3, 5, 7]  
  
>>> b = ['galicia', 'asturias', 'cantabria', 'euskadi', 'navarra']  
  
>>> b.sort()  
  
>>> b  
  
[ 'asturias', 'cantabria', 'euskadi', 'galicia', 'navarra' ]
```

Python - Collections

Lists

- `sort()` accepts a function as sorting key:

```
>>> b = [ 'caladan', 'arrakis', 'corrin', 'ix', 'giedi prime' ]  
  
>>> b.sort( key = len )  
  
>>> b  
[ 'ix', 'corrin', 'caladan', 'arrakis', 'giedi prime'  
]
```

Python - Collections

Lists

- The `bisect` module manipulates sorted lists using binary search.
- It does not check that a list is actually sorted, using it on unsorted lists yields incorrect results.

```
>>> import bisect

>>> c = [1, 2, 2, 2, 3, 4, 7]

>>> bisect.bisect( c, 2 ) # Returns insertion index

4

>>> bisect.insort( c, 6 ) # Sorted insertion

>>> c

[1, 2, 2, 2, 3, 4, 6, 7]
```

Python - Collections

Slicing

- Sections of indexed collections (such as tuples and lists) can be accessed using slice notation:

```
>>> a = [7, 2, 3, 7, 5, 6, 0, 1]
```

```
>>> a[1:5]
```

```
[2, 3, 7, 5]
```

```
>>> a[3:4] = [6, 3]
```

```
>>> a
```

```
[7, 2, 3, 6, 3, 5, 6, 0, 1]
```

Python - Collections

Slicing



- The slice `[start:stop]` includes the element in the `start` position, but not the element in the `stop` position.
- Any of them may be omitted, in which case the start or ending element of the collection is used by default:

```
>>> a[:5]
```

```
[7, 2, 3, 6, 3]
```

```
>>> a[3:]
```

```
[6, 3, 5, 6, 0, 1]
```

Python - Collections

Slicing

- Negative indices refer to the end of the array:

```
>>> a[-4:]
```

```
[5, 6, 0, 1]
```

```
>>> a[-6:-2]
```

```
[6, 3, 5, 6]
```

Python - Collections

Slicing

- The step can be modified using [start:stop:step]:

```
>>> a[::2]
```

```
[7, 3, 3, 6, 1]
```

- A particular use of the step is the reverse a sequence:

```
>>> a[::-1]
```

```
[1, 0, 6, 5, 3, 6, 3, 2, 7]
```

Python - Collections Manipulation

- Oftentimes we want to iterate the elements of a collection and their index at the same time:

```
i = 0

for x in collection:

    # do something with x, i

    i += 1
```

- This is equivalent to

```
for i, x in enumerate(collection):

    # do something with x, i
```

Python - Collections Manipulation

- `sorted()` returns a sorted list containing the elements of a collection:

```
>>> sorted( [7, 1, 2, 6, 0, 3, 2] )
```

```
[0, 1, 2, 2, 3, 6, 7]
```

```
>>> sorted( 'test string' )
```

```
[' ', 'e', 'g', 'i', 'n', 'r', 's', 's', 't', 't',  
't']
```

```
>>> sorted( set( 'test string' ) )
```

```
[' ', 'e', 'g', 'i', 'n', 'r', 's', 't']
```

Python - Collections Manipulation

- `zip()` groups the elements in several sequences into a list of tuples:

```
>>> seq1 = ['caladan', 'kaitain', 'giedi prime', 'arrakis']  
  
>>> seq2 = ['atreides', 'corrino', 'harkonnen']  
  
>>> zip( seq1, seq2 )  
  
[('caladan', 'atreides'), ('kaitain', 'corrino'),  
 ('giedi prime', 'harkonnen')]
```

- The length of the resulting list is given by the length of the shortest input sequence.

Python - Collections Manipulation

- `zip()` is commonly used to simultaneously iterate several sequences:

```
for i, (a,b) in enumerate( zip( seq1, seq2 ) ):  
    # do something with i, a, b
```

Python - Collections Manipulation

- `zip()` can be used to “unzip”:

```
>>> l = [ ('caladan', 'atreides'), ('kaitain', 'corrino'),  
        ('giedi prime', 'harkonnen') ]  
  
>>> unzip1, unzip2 = zip(*l)  
  
>>> unzip1  
  
('caladan', 'kaitain', 'giedi prime')  
  
>>> unzip2  
  
('atreides', 'corrino', 'harkonnen')
```

Python - Collections Manipulation

- `reversed()` builds an iterator over the elements of a sequence in reverse order:

```
>>> reversed( range(10) )
```

```
<listreverseiterator object at 0x7ff664159510>
```

```
>>> list( reversed( range(10) ) )
```

```
[9, 8, 7, 6, 5, 4, 3, 2, 1, 0]
```

Python - Collections

Dictionaries

- A dictionary ('dict' type) is an associative array (hashtable).
- It is similar to an array, but indexes an object using a key instead of an integer.

```
>>> empty_dict = {}
```

```
>>> d1 = { 'a': 'a value', 'b': [1,2,3,4] }
```

```
>>> d1
```

```
{'a': 'a value', 'b': [1,2,3,4]}
```

Python - Collections

Dictionaries

- The elements of a dictionary can be accessed and inserted using the same syntax as for lists and tuples:

```
>>> d1[7] = 'an integer'
```

```
>>> d1
```

```
{'a': 'a value', 'b': [1,2,3,4], 7: 'an integer'}
```

```
>>> d1['b']
```

```
[1,2,3,4]
```

Python - Collections

Dictionaries

- It can be checked whether a key is contained in a dictionary using `in`:

```
>>> 'b' in d1
```

```
True
```

```
>>> 42 in d1
```

```
False
```

Python - Collections

Dictionaries

- To remove values from a dictionary either `del` or `pop()` can be used:

```
>>> del d1['a']

>>> d1

{'b': [1,2,3,4], 7: 'an integer'}

>>> d1.pop('b')

[1,2,3,4]

>>> d1

{7: 'an integer'}
```

Python - Collections

Dictionaries

- `keys()` and `values()` return the keys and values stored in the dictionary:

```
>>> d1[5] = 'another integer'
```

```
>>> d1.keys()
```

```
[5, 7]
```

```
>>> d1.values()
```

```
['another integer', 'an integer']
```

- The keys and values are not returned in any particular order, but the orders of both lists are consistent (i.e. `d1[d1.keys()[x]] = d1.values()[x]`).

Python - Collections

Dictionaries

- Two dictionaries can be fused using `update()`:

```
>>> d1.update( { 'b': 'caladan', 'c': 'arrakis' } )  
  
>>> d1  
  
{'c': 'arrakis', 'b': 'caladan', 5: 'another integer',  
7: 'an integer' }
```

Python - Collections

Dictionaries

- A dictionary can be created from two lists:

```
mapping = {}
```

```
for key, value in zip( key_list, value_list ):
```

```
    mapping[key] = value
```

- Or:

```
>>> mapping = dict( zip( key_list, value_list ) )
```

Python - Collections

Dictionaries

- In order to be usable as a dictionary key, a Python object must be “hashable”.
- Basic types in Python are hashable, but not mutable containers (e.g., lists).
- A hashable object implements `__hash__()`, `__eq__()`, and `__cmp__()` such that:
 1. The return value of `__hash__()` does not change during the life of the object.
 2. If two objects are equal according to `__eq__()` they must share the same `__hash__()` value.
 3. `__cmp__()` must compare objects consistently.

Python - Collections

Set

- A set is an unsorted collection of unique elements:

```
>>> set( [2,2,2,1,3,3] )
```

```
set([1, 2, 3])
```

```
>>> {2, 2, 2, 1, 3, 3}
```

```
set([1, 2, 3])
```

Python - Collections

Set



Method	Alternate syntax	Description
<code>a.add(x)</code>	--	Add <code>x</code> to set.
<code>a.remove(x)</code>	--	Remove <code>x</code> from set.
<code>a.union(b)</code>	<code>a b</code>	Union of <code>a</code> and <code>b</code> .
<code>a.intersection(b)</code>	<code>a & b</code>	Intersection of <code>a</code> and <code>b</code> .
<code>a.difference(b)</code>	<code>a - b</code>	Set difference.
<code>a.symmetric_difference(b)</code>	<code>a ^ b</code>	Symmetric set difference.
<code>a.issubset(b)</code>	--	True if <code>b</code> is a subset of <code>a</code> .
<code>a.issuperset(b)</code>	--	True if <code>a</code> is a superset of <code>b</code> .
<code>a.isdisjoint(b)</code>	--	True if <code>a</code> and <code>b</code> are disjoint.

Python - Collections

Comprehensions of lists, sets, and dictionaries

- Comprehensions are “syntactic sugar” to generate new collections by operating and filtering preexisting ones.
- The basic syntax is as follows:

```
>>> [expr for val in collection if condition]
```

equals:

```
new_list = []
for val in collection:
    if condition:
        new_list.append( expr )
```

Python - Collections

Comprehensions of lists, sets, and dictionaries

- The filtering condition may be omitted:

```
>>> strings = ['a', 'an', 'the', 'cat', 'car', 'pigeon']

>>> [x.upper() for x in strings if len(x) > 2]

['THE', 'CAT', 'CAR', 'PIGEON']

>>> [x.upper() for x in strings]

['A', 'AN', 'THE', 'CAT', 'CAR', 'PIGEON']
```

Python - Collections

Comprehensions of lists, sets, and dictionaries

- Using a similar syntax we can write comprehensions of sets and dictionaries:

```
{ key-expr: val-expr for value in collection if condition }
```

```
{ set-expr for value in collection if condition }
```

Python - Collections

Comprehensions of lists, sets, and dictionaries

- We can write nested loops in a comprehension:

```
>>> tuples = ((1, 2, 3), (4, 5, 6), (7, 8, 9))

>>> [x for tup in tuples for x in tup if x > 3]

[4, 5, 6, 7, 8, 9]
```

- The order of the loops in a comprehension is the same as in an equivalent code:

```
>>> list = []

>>> for tup in tuples:
...     for x in tup:
...         if x > 3: list.append( x )
```

Python - Collections

Comprehensions of lists, sets, and dictionaries

- It is also valid to nest comprehensions:

```
>>> [[x for x in tup] for tup in tuples]  
[[1, 2, 3], [4, 5, 6], [7, 8, 9]]
```

Python - Functions



- Functions are defined using the reserved word `def`.
- `return` is used to return control to the caller and pass the result.

```
def my_func( x, y, z = 1.5 ):  
    if z > 1:  
        return z * (x+y)  
  
    else:  
        return z / (x+y)
```

Python - Functions



- If the end of the function code is reached without executing any `return` instruction, `None` is returned automatically.
- Functions receive two types of parameters: positional and keyword.
- Keyword arguments are commonly used to provide default values to optional parameters.
- In the example, `x` and `y` are positional parameters, while `z` is a keyword parameter. The function can be called in different ways:

```
>>> my_func( 5, 6,  
           z=0.7 )  
6.14
```

```
>>> my_func( 3.14, 7,  
           3.5 )  
17.99000000000002
```



Python - Functions

Name spaces, scopes, and local functions

- A function can access variables in two different scopes: global and local.
- Alternatively, variables can be explicitly defined inside a namespace.
- By default, variables assigned inside a function belong to the local scope.
- The local scope of a function is created when it is called, and initially contains its parameters.
- The local scope is destroyed when the function returns (except for closures, which we will briefly cover later).

Python - Functions

Name spaces, scopes, and local functions

```
def func():  
  
    a = []  
  
    for i in range(5):  
  
        a.append( i )
```

- When `func()` is called, `a` is created. Then, the loop is executed and 5 integers are appended to `a`. Finally, the end of the function body is reached and `a` is destroyed.



Python - Functions

Name spaces, scopes, and local functions

```
>>> a = [ ]  
  
>>> def func():  
...     for i in range(5):  
...         a.append( i )  
  
>>> func()  
  
>>> a  
[0, 1, 2, 3, 4]
```

Python - Functions

Name spaces, scopes, and local functions

```
>>> a = [ ]  
  
>>> def func():  
...     a = range(5)  
  
>>> func()  
  
>>> a  
[ ]
```

Python - Functions

Name spaces, scopes, and local functions

```
>>> a = []  
  
>>> def func():  
...     global a  
...     a = range(5)  
  
>>> func()  
  
>>> a  
[0, 1, 2, 3, 4]
```

Python - Functions

Name spaces, scopes, and local functions

- New functions can be declared anywhere in the code.
- In particular, it is legal to declare a function nested inside another function. These are called local functions, and are created when the enclosing function is called:

```
def outer_f( x, y, z ):  
    def inner_f( a, b, c ):  
        pass  
  
    pass
```

- `inner_f()` does not exist until the call to `outer_f()`. As soon as `outer_f()` returns, `inner_f()` is destroyed.
- `inner_f()` can access the variables and functions in the local scope of `outer_f()`, but it cannot add anything to it.

Python - Functions

Returning multiple values

- A notable difference w.r.t. other languages such as C/C++/Java is the ability of a function to return multiple values.

```
def f():  
  
    a = 5; b = 6; c = 7;  
  
    return a, b, c  
  
x, y, z = f()
```

- The implementation is very simple: `f()` is actually returning a single value, which happens to be a tuple.

```
>>> f()  
  
(5, 6, 7)
```

Python - Functions

Functions as objects

- A Python function is just a special type of object which defines the () operator.
- As such, it is possible to use functors (pointers to function objects) to code complex operations in a simple way.
- E.g., we want to perform cleanup operations on the strings in the following array:

```
>>> planets = [ '      Caladan ', 'Ix!', 'Ix', 'ix',
'aRraKIs', 'giedi prime##', 'Salusa secundus?' ]
```

- We want to build a list of uniform strings for its analysis. We need to apply removal of unnecessary spaces and symbols, and to fix capitalization.

Python - Functions

Functions as objects

```
import re # Regular Expression module

def clean_strings( strings ):

    result = []

    for value in strings:

        value = value.strip()

        value = re.sub('![#?]', '', value) # removes !#? symbols

        value = value.title()

        result.append( value )

    return result
```

Python - Functions

Functions as objects

```
def clean_strings_v2( strings, ops ):  
  
    result = []  
  
    for value in strings:  
  
        for function in ops:  
  
            value = function( value )  
  
        result.append( value )  
  
    return result  
  
  
def remove_punctuation( value ):  
  
    return re.sub( "[!#?]", "", value )  
  
  
clean_ops = [ str.strip, remove_punctuation, str.title ]
```

Python - Functions

Functions as objects

- It is possible to pass functions as arguments to other functions.
- E.g., Python provides mechanisms to apply a function to a list of objects:

```
>>> map( str.strip, planets )  
['Caladan', 'Ix!', 'Ix', 'ix', 'aRraKIs', 'giedi prime##',  
'Salusa secundus?']
```

Python - Functions

Anonymous (λ) functions

- An anonymous function, or lambda function, is a single-instruction functional expression the result of which is its return value:

```
def short_function(x):  
    return x*2
```

```
equiv_anon = lambda x: x*2
```

Python - Functions

Anonymous (λ) functions

- Oftentimes it is simpler to use a reference to a lambda function than to write an ad-hoc, named one:

```
>>> a = [4, 0, 1, 5, 6]
```

```
>>> map( lambda x: x*x, a )
```

```
[8, 0, 2, 10, 12]
```

```
>>> strings = ['caladan', 'ix', 'corrin', 'giedi prime']
```

```
>>> strings.sort( key = lambda x: len(set(list(x))) )
```

```
# Sort by number of different letters in the string
```

```
>>> strings
```

```
['ix', 'caladan', 'corrin', 'giedi prime']
```

Python - Functions Closures

- A closure is a dynamically-generated function which is returned by another function.
- The distinguishing characteristic of a closure is that it is capable of accessing the local scope of its generating function after the latter returns.

Python - Functions Closures

```
>>> def make_closure( a ):  
...     def closure():  
...         print( "Variable in the local scope: %d" % a )  
...     return closure  
  
>>> closure = make_closure(5)  
  
>>> closure()  
'Variable in the local scope: 5'
```

Python - Functions Closures

- In the example, a closure was created with an immutable internal state (the integer `a`).
- Mutable variables can also be used in a closure. These can be modified, dynamically altering the behavior of the closure.



Python - Functions Closures

```
>>> def make_watcher():
...     have_seen = set([])
...     def has_been_seen( x ):
...         if x in have_seen: return True
...         else: have_seen.add(x)
...         return False
...     return has_been_seen

>>> watcher = make_watcher()

>>> vals = [5, 6, 1, 5, 1, 6, 3, 5]

>>> [watcher(x) for x in vals]
[False, False, False, True, True, True, False, True]
```

Python - Functions Closures

- The local variables of a closure can be **modified**.
- No new variables can be **added** to the scope of a closure. A workaround is to add key/value pairs to a dictionary in the scope.
- Closures allow to build generic functions with plenty of options, that can be dynamically instantiated into specialized, efficient and simple variants.

Python - Functions Closures



```
>>> def format_and_pad( template, space ):  
...     def formatter( x ):  
...         return (template % x).rjust( space )  
...     return formatter  
  
>>> fmt = format_and_pad( "%.4f", 15 )  
  
>>> fmt(1.756)  
'1.7560'
```

Python - Functions

Extended syntax: *args, **kwargs

- Functions are called using a mix of positional and keyword parameters:

```
>>> func( a, b, c, d = d_value, e = e_value )
```

- Internally, this function:
 1. Receives an `args` tuple containing its positional parameters.
 2. Receives a `kwargs` dictionary containing its keyword parameters:
 3. Performs the following assignment:

```
>>> (a, b, c) = args
>>> d = kwargs.get( 'd', d_default_value )
>>> e = kwargs.get( 'e', e_default_value )
```

Python - Functions

Extended syntax: *args, **kwargs

```
def g( x, y, z=1 ): return (x+y) / z

def hello_world_then_call( f, *args, **kwargs ):
    print 'args is', args
    print 'kwargs is', kwargs
    print "Hello world! Now I'm going to call %s" % f
    return f( *args, **kwargs )

>>> hello_world_then_call( g, 1, 2, z=5 )
args is (1, 2)
kwargs is {'z': 5.0}
Hello world! Now I'm going to call <function g at 0x2dd5cf8>
```

Python - Functions

Partial function application

- Partial function application consists in creating new functions from preexisting ones by fixing some of their parameters:

```
>>> def add(x, y): return x+y
```

```
>>> add_5 = lambda y: add(5,y)
```

- Alternatively:

```
>>> from functools import partial
```

```
>>> add_5 = partial( add, 5 )
```

Python - Functions

Generators

- A generator is a function which returns a sequence of values in a lazy way, stopping its execution after each value in the sequence.
- Generators are useful to generate large, iterable sequences in a memory-efficient way (e.g., `range()` vs `xrange()`).
- Generators are declared as a function which returns a value using `yield` instead of `return`.

Python - Functions Generators

```
>>> def squares( n = 10 ):  
...     for i in xrange( 1, n+1 ):  
...         print "Generating squares from 1 to  
%d"% (n**2)  
...         yield i ** 2  
  
>>> gen = squares()  
  
>>> gen  
  
<generator object squares at 0x7fd9e3e796e0>
```

Python - Functions Generators

- When the generator function is called, no code is executed.
- Each element must be explicitly requested:

```
>>> for x in gen:
```

```
...     print x
```

```
Generating squares from 1 to 100
```

```
1
```

```
Generating squares from 1 to 100
```

```
4
```

```
.
```

```
.
```

```
.
```

Python - Functions Generators

```
def make_change( amount, coins=[1, 2, 5, 10, 20, 50], hand=[] ):

    if amount == 0: yield hand

    for coin in coins:

        if coin > amount or (len(hand) > 0 and hand[-1] < coin):
            continue

        for result in make_change(amount-coin, coins=coins,
hand=hand+[coin]):

            yield result

>>> len(list(make_change(53)))
```

Python - Functions

Generator expressions

- A compact way to declare a simple generator is to use a generator expression:

```
>>> gen = ( x ** 2 for x in xrange(100) )  
  
>>> gen  
  
<generator object <genexpr> at 0x7fd9e3e798c0>
```

- This expression is equivalent to:

```
def gen():  
  
    for x in xrange(100):  
  
        yield x**2
```

Files and Operating System

- To open a file, `open()` is called passing a relative or absolute path:

```
>>> f = open('folder/file.txt')
```

- By default, the file is opened in read-only mode ('r'). The file can be seen as a collection of lines, and iterated using a `for` loop:

```
>>> for line in f:  
...     # do something with line
```

Python Files and Operating System

Opening mode	Description
r	Read only.
w	Write only. Creates a new file, overwriting any previous one.
a	Concatenate to a file (created if it does not exist).
r+	Read-write.
b	Binary mode (usage example: 'rb').
U	Use universal end-of-line mode. Translates any end-of-line marker in the file to '\n'.

Python Files and Operating System

- To write to a file we can use the methods `write()` or `writelines()`:
- E.g., to remove white lines from a file:

```
>>> f_in = open( path, 'r' )

>>> f_out = open( 'tmp.txt', 'w' )

>>> f_out.writelines( [x for x in f_in if len(x) > 1]
)
```

Python Files and Operating System



Method	Description
<code>read([size])</code>	Reads data from the file as a string. The optional argument <code>[size]</code> is the number of bytes to read. Without it, the entire file is read.
<code>readlines([size])</code>	Same as <code>read()</code> , but returns a list of strings (one per line in the file). Without <code>[size]</code> the entire file is read.
<code>write(str)</code>	Writes <code>str</code> to the file.
<code>writelines(str)</code>	Writes a list of strings to the file, one per line.

Python Files and Operating System



Method	Description
close()	Closes the file.
flush()	Synchronizes the I/O buffer to disk.
seek(pos)	Moves the file pointer to pos.
tell()	Returns the current position of the file pointer.
closed	True if the file is closed.

NumPy



- NumPy (*Numerical Python*) is a fundamental package that enables efficient array and vector processing.
- It provides:
 - An `ndarray` class, which represents a multidimensional array and provides *vectorized* arithmetic operations and *broadcasting* capabilities.
 - *Vectorized* operations are applied to whole arrays (without using Python loops).
 - Operations for lineal algebra, randomness, signal processing, ...
 - Tools for the effective integration of Python and C/C++/Fortran.

NumPy - ndarray

002-Numpy.ipynb

- n-dimensional array object.
- Allows to perform operations on large data blocks with scalar syntax.
- Unlike lists, a NumPy array is usually homogeneous: it holds objects of a specific, predefined type.
- An ndarray contains, among others, the following two properties:
 - shape: tuple containing the array dimensions.
 - dtype: instance of the dtype class which specifies element datatype.

NumPy - ndarray Creation

002-Numpy.ipynb

Function

Description

array()	Builds a new array from the input sequence (e.g., a list). Copies input data to the new array.
asarray()	Converts the input sequence to an array, but it does not copy data if the input is already an ndarray.
arange()	Same as range() but returning an ndarray instead of a list.
ones() ones_like()	Build an array containing all 1s from a tuple specifying the desired dimensions, or copying shape and dtype from another array.

NumPy - ndarray Creation

002-Numpy.ipynb

Function	Description
<code>zeros()</code> <code>zeros_like()</code>	Build an array containing all 0s from a tuple specifying the desired dimensions, or copying shape and dtype from another array.
<code>empty()</code> <code>empty_like()</code>	Build an array without initializing its data from a tuple specifying the desired dimensions, or copying shape and dtype from another array.
<code>eye()</code> , <code>identity()</code>	Build an array with 1s in its main diagonal and 0s elsewhere. <code>identity()</code> requires square dimensions.

NumPy - ndarray

Datatypes

Type	Code	Description
int8, uint8	i8, u8	Signed/unsigned integer, 8 bits (1 byte).
int16, uint16	i16, u16	Signed/unsigned integer, 16 bits (2 bytes).
int32, uint32	i32, u32	Signed/unsigned integer, 32 bits (4 bytes).
int64, uint64	i64, u64	Signed/unsigned integer, 64 bits (8 bytes).

NumPy - ndarray

Datatypes



Type	Code	Description
float16	f2	Floating point, half precision (2 bytes).
float32	f4 o f	Floating point, single precision (4 bytes). Compatible with a C float.
float64	f64 o d	Floating point, double precision (8 bytes). Compatible with a C double and with a Python float.
float128	f16 o g	Floating point, extended precision (16 bytes).

NumPy - ndarray

Datatypes

Type	Code	Description
complex64	c8	Complex number represented as 2 x float32.
complex128	c16	Complex number represented as 2 x float64.
complex256	c32	Complex number represented as 2 x float128.

NumPy - ndarray

Datatypes

Type	Code	Description
bool	?	Boolean type (True or False).
object	o	Python object (reference).
string_	S_	String type with fixed length (1 byte per character). E.g., S10 represents a 10-character string.
unicode_	U_	Unicode type with fixed length (the number of bytes per character varies per platform). Similar to string_.

NumPy - ndarray Datatypes

002-NumPy.ipynb

- ndarray can be cast to a different datatype using astype().
- The call to astype() always creates a new copy of the array memory, even if the data is cast to its current type.
- If the cast fails for any reason, a TypeError exception is raised.

NumPy - ndarray

Broadcasting

- When two arrays are operated, NumPy must decide what operation to perform in the first place.
- If both arrays have the same dimensionality, the operation is performed in an element-wise fashion.
- Otherwise, NumPy tries to *broadcast* (replicate) the “smallest” array to match it to the largest one.

NumPy - ndarray

Broadcasting

- The broadcasting process begins by the innermost dimension (the rightmost ones), and moves towards the outer ones (to the left).
- Two dimensions are compatible if:
 1. Both have the same number of elements, or
 2. One of them has a single element.
- If none of the previous conditions is met, the operation raises a `ValueError`.
- If the second condition is met, the array with a single element in that particular dimension is copied `n` times to match both.

NumPy - ndarray

Broadcasting

002-NumPy.ipynb

- The arrays are not required to have the same number of dimensions.
Non-existing dimensions are assumed to have a single element.
- For example:

A : 256 x 256 x 3 A : 256 x 256 x 3

B : 3 B : 256 x 1 x 3

A*B: 256 x 256 x 3 A*B: 256 x 256 x 3

NumPy - ndarray

Basic indexing

002-NumPy.ipynb

- 1-dimensional arrays behave in a similar way to lists.
- An important difference is that in NumPy a slice is not a copy, but a *view* over the array data. This improves performance, but also means that slice operations may have collateral effects on the original array.
- It is possible to obtain a copy of a slice (or any array) by calling the `copy()` method of `ndarray`.

NumPy - ndarray

Basic indexing

002-NumPy.ipynb

The elements in each index of an n-dimensional array are not scalars, but (n-1)-dimensional arrays.

		axis 1		
		0	1	2
		0, 0	0, 1	0, 2
		1, 0	1, 1	1, 2
		2, 0	2, 1	2, 2





axis 0

NumPy - ndarray

Indexing with slices

002-NumPy.ipynb

- The slice syntax is also valid for indexing arrays.
- Working with n-dimensional arrays, it is possible to use any combination of basic indexing and slicing.



NumPy - ndarray

Boolean indexing

002-NumPy.ipynb

- It is possible to index an array using a boolean array.
- The boolean array must have the same dimensionality as the axis it must index.
- Boolean indexing can be used to select elements in an array using data semantics (e.g., the elements which are greater than a particular number). The conditions do not need to refer to the array being indexed.
- Boolean indexing always returns a copy of the data (not a view).
- When accessing an n-dimensional array, boolean indexing can be mixed together with basic and sliced indexing.

NumPy - ndarray

Fancy indexing

002-NumPy.ipynb

- The term *fancy indexing* describes the indexing of an array using an integer collection containing the indices of the elements to select.
- If a list of lists is provided, the behavior changes: a 1D array is returned, containing the data indexed by the tuples resulting from applying `zip()` to the input lists.
- Fancy indexing always returns a copy of the data, not a view.
- When accessing an n-dimensional array, fancy indexing can be mixed together with basic, sliced, and boolean indexing.

NumPy - ndarray

Transposition

002-NumPy.ipynb

- ndarray provides the transpose() method and the T attribute.
- For n-dimensional arrays, transpose() accepts as input a tuple specifying an arbitrary dimensional permutation.
- The T attribute is a quick way of accessing “clasic” transposition, in which a matrix is flipped over its main diagonal.
- The swapaxes() is used to swap any two axes.
- All these operations return a view over the original array.

NumPy - Universal functions

002-NumPy.ipynb

- A universal function, or `ufunc`, performs element-wise operations over an `ndarray`.
- The term describes a wrapper over a simple function which reads one or more scalars and returns a single scalar.
- Many `ufuncs` take a single element as input (e.g., `sqrt()` or `exp()`). These are *unary* `ufuncs`.
- Other `ufuncs` take two elements and return one (*binary* `ufuncs`, e.g., `add()` or `maximum()`).
- Some `ufuncs` return more than one array (e.g., `modf()` returns the integral and fractional parts of a floating point array).

NumPy - Universal functions

Unary ufuncs



Function	Operation
abs, fabs	Absolute value of integers, floating point or complex numbers. <code>fabs()</code> is a faster version for non-complex numbers.
sqrt	Square root: <code>arr ** 0.5.</code>
square	Square: <code>arr ** 2.</code>
exp	Exponentiation: <code>np.e ** arr.</code>
log, log10, log2, log1p	Natural, base 10, and base 2 logarithms. <code>log1p(x)</code> returns $\log(1+x)$.

NumPy - Universal functions

Unary ufuncs



Function	Return value
sign	The sign of each element in the array: 1 (positive), 0 (zero), -1 (negative).
ceil / floor / rint	Round up / down / to closer integer.
modf	Integral and fractional parts of an array.
isnan	True if an element is np.nan (NaN).
isfinite, isinf	Boolean array indicating if an element is finite / infinite.

NumPy - Universal functions

Unary ufuncs



Function	Return value
<code>cos, cosh, sin, sinh, tan, tanh</code>	Trigonometric and hyperbolic functions.
<code>arccos, arccosh, arcsin, arcsinh, arctan, arctanh</code>	Inverse trigonometric and hyperbolic functions.
<code>logical_not</code>	Element-wise boolean NOT.

NumPy - Universal functions

Binary ufuncs



Function	Return value
add	Addition of two input arrays.
subtract	Subtracts the second parameter from the first.
multiply	Multiplication of two input arrays.
divide / floor_divide	Division / integral division of the input.
power	Exponentiation of the bases in the first array to the exponents in the second array.

NumPy - Universal functions

Binary ufuncs



Function	Return value
maximum / fmax	Element-wise maximum. fmax() ignores NaN values.
minimum / fmin	Element-wise minimum. fmin() ignores NaN values.
mod	Element-wise modulo.
copysign	Copies the signs of the element in the second array to the values in the first.

NumPy - Universal functions

Binary ufuncs



Function	Return value
greater / greater_equal / less / less_equal / equal / not_equal	Boolean operations comparing whether each element in the first array is greater / greater or equal / less / less or equal / equal / different than each element in the second array.
logical_and / logical_or / logical_xor	Element-wise boolean operations AND / OR / XOR.

NumPy - Data processing

002-NumPy.ipynb

- Using NumPy we can express data processing algorithms that would otherwise require loops using array operations.
- This *vectorization* process usually comes together with significant performance improvements.
- We will explore different ways of using NumPy to process data.

NumPy - Data processing

Conditional logic

002-NumPy.ipynb

- `numpy.where()` allows to write a vectorized version of the expression `x if c else y`.
- Using pure Python, this can be written as:

```
>>> result = [(x if c else y) for x, y, c in zip(xarr, yarr, cond)]
```

- However, this version...
 - Is very slow.
 - works with 1-dimensional arrays only.
- The equivalent `where()` version is written:

```
>>> result = np.where( cond, xarr, yarr )
```

NumPy - Data processing

Mathematics and statistics



Method	Description
sum	Summation of all the elements of an array along a given optional axis.
mean	Arithmetic mean.
std, var	Standard deviation and variance.
min, max	Minimum and maximum.

NumPy - Data processing Mathematics and statistics



002-NumPy.ipynb

Method	Description
argmin, argmax	Indices of minimum and maximum elements.
cumsum	Cumulative sum.
cumprod	Cumulative product.

NumPy - Data processing Boolean arrays

002-NumPy.ipynb

- When working with boolean arrays, the previous methods use `True == 1` and `False == 0`. This allows to reduce boolean values using `cumsum()` / `cumprod()` instead of AND / OR.
- There are however ad-hoc methods for reducing boolean arrays: `any()` and `all()`.
- These methods work with non-boolean arrays by interpreting `non-zero == True`, `zero == False`.

NumPy - Data processing

Sorting

002-NumPy.ipynb

- As with lists, NumPy arrays can be sorted in place using `sort()`.
- `sort()` accepts an optional `axis` parameter which indicates the axis to sort over in multidimensional arrays.
- `numpy.sort()` returns a sorted copy of the array.

NumPy - Data processing

Set logic



Function	Description
<code>unique(x)</code>	Computes the unique elements in array x (1D).
<code>intersect1d(x, y)</code>	Common elements in x and y .
<code>union1d(x, y)</code>	Union of x and y .
<code>in1d(x, y)</code>	Boolean array indicating whether each element of x is in y .
<code>setdiff1d(x, y)</code>	Set difference, $x-y$.
<code>setxor1d(x, y)</code>	Symmetric set difference.

NumPy - Array I/O

Storing arrays to disk

- `np.save()` and `np.load()` store arrays to disk in binary format.
- By default, the data is stored uncompressed with extension `.npy`.

```
>>> np.save( path_string, arr )
```

.

.

.

```
>>> arr = np.load( path_string )
```

NumPy - Array I/O

Storing arrays to disk

- We can store several arrays to the same file compressed using ZIP and with .npz extension using np.savez ():

```
>>> np.savez( path_string, a = arr1, b = arr2 )
```

- Reading a .npz file NumPy returns a dictionary object that loads individual arrays in a lazy way:

```
>>> arch = np.load( path_string )
```

```
>>> arch[ 'b' ]
```

```
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

NumPy - Array I/O

Storing arrays to disk

- It is often useful to read arrays stored in text format.
- NumPy provides `np.loadtxt()` and `np.genfromtxt()`.
- Both can be adapted to a specific input format, varying comments format, delimiters, conversion functions, rows to be skipped, or columns to parse.
- `np.savetxt()` performs the reverse operation: it writes an array to a CSV-like file.
- `np.genfromtxt()` is similar to `np.loadtxt()`, but it supports treatment of missing values and structuring of the output array.

NumPy - Linear algebra

(numpy.linalg)

Function	Description
diag	Returns the elements in the diagonal of a square matrix as a 1D array, or builds a square matrix from a 1D array.
dot	Matrix product.
trace	Sum of the elements in the matrix diagonal.
det	Determinant.
eig	Eigenvalues and eigenvectors.

NumPy - Linear algebra (numpy.linalg)

Function	Description
inv	Inverse of an square matrix.
pinv	Moore-Penrose pseudo-inverse.
qr	QR decomposition.
svd	Singular-value decomposition.
solve	Solves $Ax = b$, with A a square matrix.
lstsq	Least squares solution to $y = Xb$.

NumPy - Random numbers (numpy.random)

Function	Description
seed	Changes the generator seed.
permutation	Random permutation of a sequence.
shuffle	Random permutation of a sequence (in-place).
rand	Samples a uniform distribution.
randint	Samples integers from a uniform distribution.

NumPy - Random numbers (numpy.random)

Function	Description
randn	Samples a standard normal distribution.
binomial	Samples a binomial distribution.
normal	Samples a normal distribution.
beta	Samples a beta distribution.

NumPy - Random numbers (numpy.random)

002-NumPy.ipynb

Function	Description
chisquare	Samples a χ^2 distribution.
gamma	Samples a gamma distribution.
uniform	Samples a uniform [0, 1) distribution.

Matplotlib



- The previous examples included some plotting, but without going into details.
- Matplotlib provides mostly 2D plotting capabilities, although it includes limited 3D functionalities.
- The basic supported plot types are lines, bar charts, histograms, pie charts, and variations of them.
- This section briefly introduces Matplotlib. It has a vast amount of options that cannot be covered in detail. We will later focus on higher level libraries instead.

Matplotlib - Introduction



- The `%matplotlib` magic command automatically configures iPython to show plots.
- By default, iPython detects the proper backend for the current window manager. Using `%matplotlib inline` the plots will be shown embedded inside a QT or Notebook environment.
- A basic Matplotlib plot includes the following elements:
 - `x` and `y` axes: horizontal and vertical axes, respectively.
 - Tick marks in the `x` and `y` axes.
 - Tick labels, showing axis values.
 - Drawing area, called canvas.

Matplotlib - Introduction

003-Matplotlib.ipynb

- `plot()` is used to draw lines and marks.
- It accepts pairs of `x` and `y` sequences, which must have the same lengths, together with a string indicating how to plot the data.
- Adjacent points are joined using straight lines.
- Points and lines are drawn following the requested style.
- The function returns a list of the lines which have been added to the current figure.
- If only one sequence is passed to the function, it is assumed to contain the values of the `y` axis. Values for the horizontal axis will be automatically generated as `x = range(len(y))`.

Matplotlib - Introduction

003-Matplotlib.ipynb

Function	Description
<code>figure</code>	Creates a new figure. Accepts an integer that acts as a unique identifier for this figure. This integer can be used to programmatically change the active figure.
<code>subplot(x, y, z)</code>	Divides a figure into a mesh of subfigures with x rows and y columns. Besides, it activates subfigure number z ($1 < z < x * y$, row major order).

Matplotlib - Introduction

003-Matplotlib.ipynb

Function	Description
bar / barh	Create vertical / horizontal bar charts. To plot stacked bars the bottom parameter is used, indicating the starting point for the new bars.
boxplot	Box plots.
scatter	Draws points, but it does not join them with lines unlike plot().
hist	Histograms.
pie	Pie charts.

Matplotlib - Design

003-Matplotlib.ipynb

Function	Description
title	Plot title. Like all other text-related Matplotlib functions, it accepts LaTeX syntax.
xlim / ylim	Configures the limits of the x / y axes.
autoscale	Automatic axes limits.
xticks / yticks	Configures the ticks for each axis and, optionally, the tick labels.

Matplotlib - Design

003-Matplotlib.ipynb

Function	Description
axis	Configures both axes' limits.
axhline / axvline	Draws a horizontal / vertical line,
axhspan / axvspan	Draws rectangles that cover the entire width / height of the plot.

Matplotlib - Design Axes

003-Matplotlib.ipynb

- A reference to an instance of `matplotlib.axes.Axes` allows to configure plots with a high level of detail.
- Such reference can be obtained by calling `matplotlib.pyplot.gca()` (Get Current Axes).
- E.g., calling `ax.xaxis.set_major_locator(matplotlib.ticker.MultipleLocator(10))` sets ticks in the x axis of the `ax` axes in values multiple of 10.
- E.g., calling `ax.xaxis.set_major_formatter()` allows to specify functions that will be used to format the tick labels in the x axis of the `ax` axes.

Matplotlib - Design

Legends and annotations

003-Matplotlib.ipynb

- All plotting functions (`plot()`, `hist()`, etc.) accept a `label` parameter indicating the name to use in the legend.
- `matplotlib.pyplot.legend()` automatically builds a legend from a list of handles. If no such list is provided, all the elements in the figure with a label not starting with “`_`” are listed.
- Otherwise, we can provide the lists of lines, etc. returned by each plotting function to specify which elements in the figure should be added to the legend.
- `matplotlib.pyplot.annotate()` adds a textual annotation and the specified place in the figure.

Matplotlib - Styles

003-Matplotlib.ipynb

Property	Value type	Description
alpha	float	Transparency used for the element.
color	Matplotlib color	Line / marker color.
dashes	Sequence	Line pattern.
label	string	Text to label this element in the plot legend.
linestyle	See docs	Line type.
linewidth	float	Line width, in points.

Matplotlib - Styles

003-Matplotlib.ipynb

Property	Value type	Description
marker	See docs	Marker used for points in a line.
mec	Matplotlib color	Marker Edge Color.
mew	float	Marker Edge Width.
mfc	Matplotlib color	Marker Face Color.
markersize	float	Marker size, in points.

Matplotlib - Styles

003-Matplotlib.ipynb

Property	Value type	Description
solid_capstyle	['butt' 'round' 'projecting']	End of line style.
solid_joinstyle	['miter' 'round' 'bevel']	Line union style.
visible	[True False]	Visibility.
xdata	numpy.ndarray	X axis data.
ydata	numpy.ndarray	Y axis data.
Zorder	Número	Stacking order in the Z axis.

Matplotlib - Colors

003-Matplotlib.ipynb

- The `matplotlib.colors` module includes utilities for defining and converting colors.
- Basic predefined colors can be referenced using a single letter: b (*blue*), g (*green*), r (*red*), c (*cyan*), m (*magenta*), y (*yellow*), k (*black*), w (*white*).
- Shades of gray can be codified using a floating point number in [0, 1].
- Other colors can be specified using different formats:
 - HTML hex string: "#eeeeff"
 - (R, G, B) tuple, with R, G and B in [0, 1].
 - HTML string: "red", "burlywood", "chartreuse", ...

Matplotlib - Colors Colormaps

003-Matplotlib.ipynb

- The `matplotlib.cm` module includes a set of colormaps to use with plots.
- Colormaps are useful when using `imshow()`, which is similar to `plot()` but interprets the values in the input array as indices to a colormap. If the input array is 2-dimensional, an image will be plotted.

Matplotlib - Saving to file

003-Matplotlib.ipynb

- Figures can be saved to a file using `matplotlib.pyplot.savefig()`.
- The file format to use is inferred from the extension in the provided path.
- The most relevant parameters that control the quality of the stored figure are:
 - `dpi`: dots per inch.
 - `bbox_inches`: inches of whitespace surrounding the figure.
- The figure does not need to be stored to a file: it can be written to any object which supports I/O, such as `StringIO`.

Matplotlib - 2D histograms



Parameter to `hist()`

Description

<code>bins</code>	Number of classes to use in the histogram, or sequence containing the boundaries between classes.
<code>range</code>	Range for each class. If <code>bins</code> is not used, this can be provided as a sequence.
<code>normed</code>	If <code>True</code> , values are normalized and the result is a probability density function.

Matplotlib - 2D histograms

003-Matplotlib.ipynb

Parameter to `hist()`

Description

histtype	By default, bar chart. Other values: <ul style="list-style-type: none">• barstacked: stacks bars when working with multiple data series.• step: line without filling.• stepfilled: line with filling.
align	How to align bars for each class: mid, left, or right.
color	Colors to be used.
orientation	horizontal or vertical.

Matplotlib - Pie charts

003-Matplotlib.ipynb

Parameter to `pie()`

Description

explode	Fraction of the pie ratio to use as offset for each slice.
autopct	String or function indicating how to label each slice with a numeric value.
pctdistance	Ratio between the center of each slice and the start of the <code>autopct</code> text.
labeldistance	Radial distance to draw labels.
startangle	Rotation angle of the origin.
wedgeprops	Dictionary containing slice properties.

Matplotlib - 3D plots

003-Matplotlib.ipynb

- Although Matplotlib is focused in 2D plots, there are several toolkits provided limited 3D capabilities.
- The `mpl_toolkits.mplot3d` module provides methods to create 3D scatter plots, surfaces, lines and meshes. Its interface is very similar to the original Matplotlib.
- Main difference: axes are instances of `mpl_toolkits.mplot3d.Axes3D`. Projections are performed by specialized 3D classes, while other parts of the figure (labels, ticks, etc.) are managed by vanilla Matplotlib.

Bibliography

- *Learning Python*. Mark Lutz. O'Reilly Media, 2013 (5^a edición).
- *Programming Python*. Mark Lutz. O'Reilly Media, 2011 (4^a edición).

Data Analytics: Pandas

- Intro to Pandas.
- Data I/O.
- Data wrangling.
- Visualization.
- Aggregation.
- Time series.

Pandas



- Pandas provides data structures and methods to improve the structured data processing capabilities of native Python.
- The basic data structure in Pandas is the DataFrame (similar to `data.frame` in R). It is a 2D table, conceptually similar to an Excel spreadsheet.
- Pandas combines the array processing capabilities of NumPy with the flexibility of spreadsheets and relational databases.
- Provides indexing, reshaping, splitting, aggregation, and selection of subsets of data.

Pandas



- Pandas design objective is to provide new data management capabilities to the Python ecosystem:
 - Data structures with indexed axes supporting explicit or implicit data alignment.
 - Seamless processing of time series (timestamp-indexed data).
 - Arithmetic operations and reductions along axes.
 - Flexible management of unknown (null) data.
 - Merges and similar operations typical of relational databases.

Pandas - Data structures

005-Pandas-Intro.ipynb

Series

- A Series object represents an object conceptually similar to a 1D array.
- It actually contains two different arrays: a data array and an index array, which labels data.
- Both arrays are essentially NumPy arrays with their own datatypes.
- Series are similar to dictionaries, as they can be seen as key-value pairs. In fact, a constructor is provided to build a Series object from a dict.

Pandas - Data structures

005-Pandas-Intro.ipynb

DataFrame

- A DataFrame object represents a tabular structure, similar to a spreadsheet (or a `data.frame` in R).
- It contains an ordered collection of columns, each of which may have a different datatype (numeric, string, etc.).
- Includes an index on its rows and another one on its columns. In this sense, it can be seen as an aggregation of series, all of them sharing the same index.
- Internally, the data is stored in a 2D format (bidimensional NumPy array), although higher dimensional data may be represented using hierarchical indices.

Pandas - Data structures

DataFrame() constructor

Parameter (type)	Description
2D ndarray	Data array, with optional row and column labels.
dict of arrays, lists, or tuples	Each sequence becomes a column of the DataFrame. All of them might have the same length.
dict of Series	Each value becomes a column. The indices in each series are unified if an explicit index is not provided.

Pandas - Data structures

DataFrame() constructor

Parameter (type)	Description
dict of dicts	Each internal dictionary becomes a column. The keys in the different dictionaries are unified as in a dict of Series.
list of dicts or Series	Each item becomes a row in the DataFrame. The union of the keys of the dictionaries or the indices of the series is used for the column labels.
list of lists or tuples	Same as providing an 2D ndarray.

Pandas - Data structures

DataFrame () constructor

Parameter (type)	Description
DataFrame	The indices already in the old DataFrame are used, unless different ones are explicitly provided.

Pandas - Data structures

005-Pandas-Intro.ipynb

Index

- Index objects are responsible for storing axis labels and names.
- Any other array or sequence type provided as an index when building a Series or DataFrame object is internally converted to an Index.
- Index objects are immutable. This guarantees referential integrity when shared by different structures.
- Index provides methods and attributes to support set logic and value inspection.

Pandas - Data structures

Index subclasses

Class	Description
Index	Array of generic Python objects.
Int64Index	Integers.
MultiIndex	Hierarchical index, representing multiple indexing levels in a single axis. Similar to a tuple array.
DatetimeIndex	Timestamp with nanosecond resolution.
PeriodIndex	Time periods.

Pandas - Data structures

Index methods



Method	Description
append	Concatenates additional <code>Index</code> objects, producing a new object.
diff	Set difference.
intersection	Set intersection.
union	Set union.
isin	Computes a boolean array marking whether each of the entries in an <code>Index</code> is included in another collection.
delete	Creates a new <code>Index</code> object by removing an element from the original one.

Pandas - Data structures

Index methods



Method	Description
drop	Creates a new <code>Index</code> object by removing a set of elements from the original one.
insert	Creates a new <code>Index</code> object by inserting a new element into the original one.
is_monotonic	<code>True</code> if each element is greater or equal than the previous one. Alias for <code>is_monotonic_increasing()</code> . An <code>is_monotonic_decreasing()</code> method is also provided.
is_unique	<code>True</code> if the index has no duplicate elements.
unique	Computes an array containing the unique elements in the index.

Pandas - Essential functions

005-Pandas-Intro.ipynb

- Reindexing: `reindex()` method.
- Removing entries: `drop()` method.
- Indexing, selection, and filtering: operator `[]` and attributes `DataFrame.loc` and `DataFrame.iloc`.
- Arithmetic operations: operators `+`, `-`, `*`, and `/`, and methods `add()`, `sub()`, `mul()` y `div()`.
- Functional application and mapping: basic methods (`mean()`, `sum()`, ...), method `Series.map()`, methods `DataFrame.apply()`, and `DataFrame.applymap()`.
- Sorting and classification: methods `sort_values()`, `sort_index()`, and `rank()`.
- Managing indices with duplicates: method `Index.is_unique()`.

Pandas - Essential functions

reindex () parameters

Parameter	Description
index	New sequence to use as index.
method	Interpolation method (ffill or bfill).
fill_value	Value to use as a placeholder for null data.
limit	Maximum number of elements to fill using interpolation.
level	Hierarchical level to reindex.
copy	Marks whether data should be copied in case the new and the old indices are equivalent (True by default).

Pandas - Essential functions

DataFrame indexing

Syntax	Description
obj [val]	Selects a column or subset of columns, except if val is an array or the DataFrame is boolean, in which case it filters columns.
obj .loc [val]	Selects a row or subset of rows by label.
obj .loc [:, val]	Selects a column or subset of columns, by label.
obj .loc [val1, val2]	Selects both rows and columns, by label.

Pandas - Essential functions

DataFrame indexing

Syntax	Description
<code>reindex()</code>	Reorganizes one or more axes according to new indices.
<code>xs()</code>	Returns a cross-section of the dataframe attending to its labels.
<code>icol()</code> / <code>irow()</code>	Selects a single row / column attending to its location.
<code>get_value()</code> / <code>set_value()</code>	Selects a single value attending to row/column labels.

Pandas - Descriptive statistics

Reduction parameters

Parameter	Description
axis	The axis on which to perform the reduction (rows=0, columns=1).
skipna	Whether to exclude null values. True by default.
level	Reduction grouping by level in hierarchical indices.

Pandas - Descriptive statistics

Methods

Method	Description
count	Number of non-null values.
describe	Computes several statistics of a Series or the columns of a DataFrame.
min / max	Minimum / maximum value.
argmin / argmax	Location of the minimum / maximum value.
idxmin / idxmax	Index (label) of the minimum / maximum value.
quantile	Returns the specified p-quantile.
sum	Summation.

Pandas - Descriptive statistics

Methods



Method	Description
mean	Arithmetic mean.
median	Median.
mad	Mean absolute deviation.
var	Variance.
std	Standard deviation.
skew	Skewness (third standardized moment).
kurt	Kurtosis (fourth standardized moment).

Pandas - Descriptive statistics Methods

Method	Description
cumsum	Cumulative sum.
cummin, cummax	Cumulative minimum / maximum.
cumprod	Cumulative product.
diff	First order differences.
pct_change	Percentage change.

Pandas - Unknown data

005-Pandas-Intro.ipynb

- Unknown or null data are common in most data analysis applications.
- Pandas simplifies the management of null data, e.g., all descriptive statistics functions automatically omit unknown values.
- `numpy.nan` is used as the default placeholder for unknown data.
None will also be treated as a null value by Pandas.

Pandas - Unknown data fillna() parameters

Parameter	Description
value	Escalar value or dictionary to use for filling.
method	Interpolation type (“ffill” or “bfill”).
axis	Axis to fill (by default 0, i.e., rows).
inplace	Modifies the object inplace, instead of creating a new copy.
limit	Maximum number of consecutive values to fill using interpolation.

Pandas - Hierarchical indices

005-Pandas-Intro.ipynb

- Hierarchical indexing allows to works (at a conceptual level) using DataFrame objects with more than 2 dimensions.
- It is implemented through the addition of different *levels* to the row and/or column indices of the table.
- Hierarchical indices are implemented by the `MultiIndex` class.
- It is possible to apply reduction operations to the different levels in the index hierarchy, obtaining a DataFrame (instead of a Series) after the reduction.

Pandas - Other considerations

Integer indexing

- Pandas objects with integer indices can be confusing, as the semantics of position- and label-based indexation vary.

```
obj = Series( range(5), index=['a', 'b', 'c', 'd',  
'e'] )
```

```
obj[-1]
```

```
??
```

```
obj = Series( range(5), index=range(5) )
```

```
obj[-1]
```

```
??
```

Pandas - Other considerations

Integer indexing



- For objects with an integer index, Pandas cannot decide whether the user wants to apply location- or label-based indexing.
- It solves the ambiguity by always using label-based indexing.
- Provides methods `Series.iget_value()` and `DataFrame.irow()` for positional-based indexing in integer-indexed data.

Pandas: Data I/O

Reading / writing text

006-Pandas-DataIO.ipynb

- Pandas provides a set of methods to create DataFrame objects from tabular data stored in text format. The most useful ones are `read_csv()` and `read_table()` (which are nowadays mostly equivalent), which include options for:
 - Using one or more columns in the text file as DataFrame indices.
 - Naming columns using parameters or extracting names from the file.
 - Inferring types and performing data conversions.
 - Parsing dates, including combining several columns into a single one.
 - Iterating over chunks of large files (to fit data into memory).
 - Cleaning data: ignoring some rows or columns, comments, etc.
- Automatic type inference implies that it is not required to specify column types. Management of data and other non-basic types takes extra effort.

Pandas: Data I/O

`read_csv / read_table` parameters

Parameter	Description
<code>path</code>	URL of the file to open.
<code>sep / delimiter</code>	Regular expression to use for separating fields.
<code>header</code>	Row number containing the column names (0 by default), or <code>None</code> if correlative integers should be used.
<code>index_col</code>	Columns to build the <code>DataFrame</code> index.
<code>names</code>	Names for the <code>DataFrame</code> columns. Combine with <code>header=None</code> .
<code>skiprows</code>	Number of rows to ignore at the beginning of the file, or list of row numbers to ignore.

Pandas: Data I/O

read_csv / read_table parameters

Parameter	Description
na_values	Values that should be considered marks of unknown data.
comment	Regular expression to mark the beginning of a comment.
parse_dates	Tries to parse dates to a <code>datetime</code> object. <code>False</code> by default. If <code>True</code> , Pandas tries to parse all columns as dates. Alternatively, it can be a list of specific columns to parse as dates. If an element in the list is a tuple/list, it will try to combine the specified columns to parse a single date.
keep_date_col	If several columns are used to build a single date column, do not keep the joined columns. <code>True</code> by default.

Pandas: Data I/O

read_csv / read_table parameters

Parameter	Description
converters	Dictionary containing a mapping of columns to parsing functions. Each corresponding function will be applied to all the elements of a given column and the result will be inserted into the DataFrame.
dayfirst	When parsing potentially ambiguous dates, assume dates in international format (DD/MM/YYYY). False by default.
date_parser	Function to use to parse dates.
nrows	Number of rows to read, starting from the beginning of the file.
iterator	If True, the return value will be a TextFileReader object, to process the file in a chunk-by-chunk manner.

Pandas: Data I/O

`read_csv / read_table` parameters

Parameter	Description
<code>chunksize</code>	If chunk-by-chunk processing is active, size of each chunk.
<code>skip_footer</code>	Number of rows to ignore at the end of the file.
<code>verbose</code>	If <code>True</code> , print information about the parsing process.
<code>encoding</code>	Character encoding, e.g., ‘utf-8’.
<code>squeeze</code>	If <code>True</code> , in case the result contains a single column, return a <code>Series</code> .
<code>thousands</code>	Thousands separator, e.g., ‘,’ or ‘.’

Pandas: Data I/O

Binary formats

006-Pandas-DataIO.ipynb

- The `pickle` module in the Python standard library provides a convenient method to serialize (marshall) and deserialize (unmarshall) objects to binary format.
- Pandas provides the methods `to_pickle()`/`read_pickle()` which store/read data to/from `pickle` files.
- Using `pickle` is discouraged for long term storage, since `pickle` does not guarantee backwards compatibility.
- Pandas also provides methods to read and write Excel, HDF5, Stata, and HTML files, among others.

Pandas: Data I/O

HTML and web services

006-Pandas-DataIO.ipynb

- Many web sites include APIs that provide access to data sources in JSON format.
- There are several ways to access these services from Python.
- A simple way is through the `requests` package.
- The responses to requests will be converted to a JSON object.
- It is trivial to build a Pandas object from JSON, as we have seen.

Pandas: Data I/O

Databases

006-Pandas-DataIO.ipynb

- Text and binary files are ultimately inefficient to store large amounts of data.
- Databases, both relational and non-relational, are one of the most common data sources in computer science.
- Pandas provides methods to load data from SQL queries. The `pandas.io.sql` module allows to execute SQL sentences and process the results.
- Other types of non-SQL databases, such as MongoDB, store objects in different formats: JSON, text, etc. Loading data mechanisms vary on a per-case basis.

- Much of the programming work in data analysis is spent on data preparation: loading, cleaning, transforming, and rearranging.
- Pandas and Python provide a set of flexible, high level data manipulation tools suitable for these preparation tasks:
 - Merging and combination of data: `merge()`, `concat()`, `combine_first()`.
 - Reshaping and pivoting: `stack()`, `unstack()`, `pivot()`.
 - Data transformation: removing duplicates, functional application, replacing values, renaming axes, discretization and binning, detecting and filtering outliers, permutation, and random sampling.
 - String manipulation: normalizing, cleaning, use of regular expressions...

Pandas: Data wrangling

Combining and merging datasets

007-Pandas-Data-wrangling.ipynb

- `pandas.merge()` connects rows in DataFrame based on one or more keys. It is similar to a `join` operation in a relational database.
- `pandas.concat()` glues or stacks together objects along an axis:
 - How are the new axes labeled? (union, intersection, ...)
 - Are the original groups identifiable in the resulting object?
 - Which axis should be used for concatenating?
- `combine_first()` enables splicing together overlapping data to fill in missing values in one object with values from the other.

Pandas: Data wrangling

merge () parameters



Parameter	Description
left	Left hand side operand of the merge.
right	Right hand side operand of the merge.
how	[‘inner’ ‘outer’ ‘left’ ‘right’]. ‘inner’ by default.
on	Column names to join on. Must be found in both DataFrame objects. If not specified, will use all columns with matching names.

Pandas: Data wrangling

merge () parameters



Parameter	Description
left_on	Columns in <i>left</i> operand to use as join keys.
right_on	Columns in <i>right</i> operand to use as join keys.
left_index	[True False]. Use row index in <i>left</i> as its join key.
right_index	[True False]. Use row index in <i>right</i> as its join key.

Pandas: Data wrangling

merge () parameters



Parameter	Description
sort	[True False]. Sort merged data lexicographically by join keys. True by default. Disable to get better performance in some cases on large datasets.
suffixes	Tuple of string values to append to column names in case of overlap. ('_x', '_y') by default.
copy	If False, avoid copying data into resulting data structure in some exceptional cases. By default always copies.

Pandas: Data wrangling

concat() parameters



Parameter	Description
objs	List or dictionary of Pandas objects to be concatenated. The only required argument.
axis	Axis to concatenate along. 0 by default (rows).
join	[‘inner’ ‘outer’]. ‘outer’ by default. Whether to intersect (inner) or union (outer) together indices along the other axes.
join_axes	Specific indices to use for the other (n-1) axes instead of performing union/intersection logic.

Pandas: Data wrangling

concat() parameters



Parameter	Description
keys	Values to associate with objects being concatenated, forming a hierarchical index along the concatenation axis.
levels	Specific indices to use as hierarchical index level or levels if keys passed.
names	Names for created hierarchical levels if keys and/or levels passed.
verify_integrity	Check new axis in concatenated objects for duplicates and raise exception if so. False by default.
ignore_index	Do not preserve indices along concatenation axis.

Pandas: Data wrangling

Reshaping and pivoting

007-Pandas-Data-wrangling.ipynb

- A different kind of transformations is its reshaping, sometimes called pivoting.
- It consists in transposing rows and columns, modifying data dimensionality.
- The basic reshaping and pivoting operations provided by Pandas are:
 - `stack()`: pivots from the columns in the data to the rows.
 - `unstack()`: pivots from the rows into the columns.
 - `pivot()`: reshapes rows and columns, allowing to transform tables in “long” format to “wide” format in a single step.

Pandas: Data wrangling

Data transformations

007-Pandas-Data-wrangling.ipynb

- We have focused on structural modifications. There are other types of transformations which focus on the data:
 - Removing duplicates: `duplicated()`, `drop_duplicates()`.
 - Functional application: `apply()`, `map()`, `applymap()`.
 - Value substitution: `replace()`.
 - Index renaming: `rename()`.
 - Discretization and binning: `cut()`, `qcut()`.
 - Detecting and filtering outliers.
 - Permutation and random sampling: `permutation()`, `take()`.
 - Computing indicators/dummy variables: `get_dummies()`.

Pandas: Data wrangling

String manipulation

007-Pandas-Data-wrangling.ipynb

- One of Python's most popular characteristics is its string manipulation routines.
- The `str` class provides methods for conveniently performing many string operations, such as searching, substitutions, splitting, etc.
- Many of these operations accept regular expressions as parameters.
- Pandas adds string functionality, as it allows to apply string operations over data tables automatically handling unknown values.

Pandas: Data wrangling

str methods



Method	Description
count	Returns the number of non-overlapping occurrences of substring in the string.
endswidth / startswith	Returns True if a string ends with suffix / starts with prefix.
join	Use string as a delimiter for concatenating a sequence of other strings.
index	Return position of first character in substring if found in the string. Raises ValueError if not found.

Pandas: Data wrangling

str methods



Method	Description
find	Like <code>index()</code> , but returns <code>-1</code> if not found.
rfind	Like <code>find()</code> , but returns position of last occurrence.
replace	Replace occurrences of string with another string.
strip / rstrip / lstrip	Trim whitespace, including newlines.
split	Break string into list of substrings using passed delimiter.

Pandas: Data wrangling

str methods



Method	Description
lower / upper	Convert alphabet characters to lowercase or uppercase, respectively.
ljust / rjust	Left / right justify. Pad opposite side of string with spaces (or some other fill character) to return a string with a minimum width.

Pandas: Data wrangling

Regular expressions

007-Pandas-Data-wrangling.ipynb

- Regular expressions provide a flexible way to search or match string patterns in text.
- A regular expression is a string containing certain elements, called “special characters”, with a specific semantic.
- A regular expression (*regex*) describes a pattern to match in the text, or a way to manipulate it.
- The `re` module includes functions for regular expression application. It has three types of functions: pattern matching, substitutions, and splitting.

Pandas: Data wrangling regex methods



Method	Description
.findall / .finditer	Return all non-overlapping matching patterns in a string as a list / iterator.
.match	Match pattern at start of string and optionally segment pattern components into groups. Returns a <code>match</code> object, or <code>None</code> .
.search	Scan string for match to pattern; returning a match object if so. The match can be anywhere in the string, as opposed to <code>match()</code> .
.split	Break string into pieces at each occurrence of pattern.
.sub, .subn	Replace all / first n occurrences of pattern with replacement expression.

Pandas: Data wrangling

Vectorized string functions in Pandas



Method	Description
cat	Concatenate strings element-wise with optional delimiter.
contains	Return boolean array if each string contains pattern.
count	Count occurrences of pattern.
endswith / startswith	Applies homonym functions in <code>str</code> element-wise.
findall	Compute list of all occurrences of pattern for each string.

Pandas: Data wrangling

Vectorized string functions in Pandas



Method	Description
get	Index into each element (retrieve i-th element).
join	Join strings in each element of the Series with passed separator.
len	Compute length of each string.
lower / upper	Convert cases.
match	Use <code>re.match</code> with passed regex on each element.
pad	Add whitespace to left, right, or both sides of string.

Pandas: Data wrangling

Vectorized string functions in Pandas



Method	Description
center	Equivalent to pad(side="both").
repeat	Duplicate values.
replace	Replace occurrences of pattern with some other regex.
slice	Slice each string in the Series.
split	Split strings on delimiter.
strip / rstrip / lstrip	Trim whitespace, including newlines, element-wise.

Pandas: Plotting and visualization

008-Pandas-Plotting.ipynb

- Matplotlib provides a powerful framework for plotting and visualization, but it is a low-level tool.
- Building a plot involves configuring several Python objects without implicit semantics.
- Pandas objects are a centralized storage of data, with semantics at least partially known.
- Pandas objects provide a `plot()` method which builds complex plots in a convenient manner.

Pandas: Plotting and visualization

plot() parameters

Parameter	Description
ax	Matplotlib subplot object to plot on. If nothing passed, uses active subplot.
kind	[‘line’ ‘bar’ ‘barh’ ‘hist’ ‘box’ ‘kde’ ‘density’ ‘area’ ‘pie’ ‘scatter’ ‘hexbin’]. Type of chart.
logx / logy	Use logarithmic scale on the X / Y axis.
use_index	Use the object index for tick labels.
rot	Rotation angle of tick labels.
xticks / yticks	Values to use for X / Y axis ticks.

Pandas: Plotting and visualization

plot() parameters



Parameter	Description
xlim / ylim	X / Y axis limits (specified as [min, max]).
grid	[True False]. Display axis grid.
subplots	[True False]. Plot each DataFrame column in a separate subplot.
sharex / sharey	[True False]. If subplots=True, share the same X / Y axis, linking ticks and limits.
layout	Tuple indicating the geometry of subplots to use.
figsize	Figure size.

Pandas: Plotting and visualization

plot() parameters



Parameter	Description
title	Plot title.
legend	[True False 'reverse']. Add a subplot legend.
style	Dictionary matching each column with the style to use for plotting it.
loglog	Use log scale for both axes.
fontsize	Font size to use for ticks.
colormap	Color map to index.

Pandas: Plotting and visualization

plot() parameters



Parameter	Description
colorbar	[True False]. Whether to draw the value legend for 'scatter' and 'hexbin' plot types).
table	If a Series or DataFrame object is provided, includes it in the plot. Useful to combine plots and tables.
stacked	Create stacked plot.
sort_columns	[True False]. Whether to sort columns lexicographically before plotting.
secondary_y	List of columns that should be referenced to a secondary Y axis.

Pandas: Plotting and visualization

plot() parameters




Parameter	Description
mark_right	[True False]. When a secondary axis is used, automatically add the suffix “(right)” to the legends of the series referenced to it.
**kwds	Parameters not processed by Pandas will be passed to Matplotlib.

Pandas: Data aggregation



- One of the reasons for the popularity of relational databases and SQL is the ease with which data can be joined, filtered, transformed, and aggregated.
- However, query languages like SQL have limited expressiveness. Pandas allows to implement *split-apply-combine* operations conveniently:
 - Split a Pandas object into pieces using one or more keys.
 - Compute group summary statistics.
 - Apply a varying set of functions to each column of a DataFrame.
 - Compute pivot tables and cross-tabulations.
 - Perform quantile analysis and other data-derived group analyses.

Pandas: Data aggregation



Pandas: Data aggregation

009-Pandas-Aggregation.ipynb

- A critical aspect of these transformations is how to categorize data.
- A Pandas object is *split* into groups based on one or more keys, applied on a particular axis.
- The Pandas mechanism for this operation is `groupby()`.
- A grouping key can take many forms, e.g.:
 - A list or array of values that is the same length as the axis being grouped.
 - A value indicating a column name in a `DataFrame`.
 - A dict or `Series` giving correspondences between the values on the axis being grouped and the group names.
 - A function to be invoked on the axis index or the individual labels of the index.

Pandas: Data aggregation

Aggregation functions

009-Pandas-Aggregation.ipynb

- An *aggregation functions* is any transformation which produces a scalar value from an array (also called *reduction*).
- Aggregation functions as implemented by the `GroupBy` class have been optimized and are computed on the original data of the `DataFrame` or `Series`.
- Applicable aggregation functions are not limited to this subset: any function, including user-defined functions, can be applied to a grouped dataset.

Pandas: Data aggregation

GroupBy methods



Method	Description
count	Number of non-NA values in the group.
sum	Sum of non-NA values.
mean	Mean of non-NA values.
median	Arithmetic median of non-NA values.

Pandas: Agregación

Métodos en GroupBy

Método	Descripción
std / var	Unbiased standard deviation / variance.
min / max	Minimum / maximum of non-NA values.
prod	Product of non-NA values.
first / last	First / last non-NA value.

Pandas: Group-wise operations



009-Pandas-Aggregation.ipynb

- Aggregation is only one kind of group operation: accepts functions that reduce a one-dimensional array to a scalar value.
- In the general case, we want to apply any kind of operation to grouped data.
- This is done using `transform()` and `apply()`:
 - `transform()` broadcasts the result of an aggregation over the original members of the group.
 - `apply()` applies a function to each group and combines the results using `pandas.concat()`.

Pandas: Pivot tables and cross-tabulation



- A pivot table is a data summarization tool which aggregates a table by one or more keys, arranging the data in a rectangle with some groups along rows and some along columns.
- It can be built using `groupby()`, but `pivot_table()` provides a more convenient high-level interface.
- A cross-tabulation is a special case of a pivot table that computes group frequencies.
- Could also be built manually using several functions, but `crosstab()` simplifies the process.

Pandas: Time series

010-Pandas-Time-series.ipynb

- Any dataset which includes observations at many points in time forms a time series.
- Many time series are *fixed frequency*: data points occur at regular intervals.
- Others are *irregular*: without a fixed offset between data points.
- How time series are referred depends on the application. Among others:
 - *Timestamps*: specific instants in time.
 - *Fixed periods*, such as the month of January 2007 or the full year 2010.
 - *Intervals* of time, indicated by start and end timestamps.
 - Elapsed time relative to a particular fixed start time.
- Pandas provides a standard set of time series tools and data algorithms to slice and dice, aggregate, resample, etc.

Pandas: Time series

Base frequencies

Alias	Offset type	Description
D	Day	Calendar daily.
B	BusinessDay	Business daily.
H	Hour	Hourly.
T / min	Minute	Minutely.
S	Second	Secondly.
L / ms	Milli	Millisecond.
U	Micro	Microsecond.

Pandas: Time series

Base frequencies



Alias	Offset type	Description
M	MonthEnd	Last calendar day of month.
BM	BusinessMonthEnd	Last business day of month.
MS	MonthBegin	First calendar day of month.
BMS	BusinessMonthBegin	First business day of month.
W-MON, W-TUE, ...	Week	Weekly on given day of week: MON, TUE, WED, THU, FRI, SAT, or SUN.
WOM-1MON , WOM-1TUE, ...	WeekOfMonth	Generate weekly dates on the first, second, third, or fourth week of the month. For example, WOM-3FRI for the 3rd Friday of each month.

Pandas: Time series

Base frequencies



Alias	<i>Offset type</i>	Description
Q-JAN, ...	QuarterEnd	Quarterly dates anchored on last calendar day of each month, for year ending in indicated month.
BQ-JAN, ...	BusinessQuarterEnd	Quarterly dates anchored on last business day of each month, for year ending in indicated month.
QS-JAN, ...	QuarterBegin	Quarterly dates anchored on first calendar day of each month, for year ending in indicated month.
BQS-JAN, ...	BusinessQuarterBegin n	Quarterly dates anchored on first business day of each month, for year ending in indicated month.

Pandas: Time series

Base frequencies

Alias	Offset type	Description
A-JAN, ...	YearEnd	Annual dates anchored on last calendar day of given month.
BA-JAN, ...	BusinessYearEnd	Annual dates anchored on last business day of given month.
AS-JAN, ...	YearBegin	Annual dates anchored on first calendar day of given month.
BAS-JAN , ...	BusinessYearBegin	Annual dates anchored on first business day of given month.

Pandas: Time series resample() parameters



Parameters	Description
freq	String or DateOffset indicated desired resample frequency.
how	Function name or array function producing aggregated value.
axis	Axis to resample on, default to 0 (rows).
fill_method	How to interpolate when upsampling ('ffill' or 'bfill').
closed	In downsampling, which end of each interval is closed (inclusive). Defaults to 'right'.
label	In downsampling, how to label the aggregated result, with the right or left bin edge. Defaults to 'right'.

Pandas: Time series

resample() parameters

Parameters	Description
loffset	Time adjustment to the bin labels, such as '-1s' / Second (-1) to shift the aggregate labels one second earlier.
limit	When forward or backward filling, the maximum number of periods to fill.
kind	Aggregate to periods ('period') or timestamps ('timestamp'); defaults to kind of index the time series has.
convention	When resampling periods, the convention ('start' or 'end') for converting the low frequency period to high frequency. Defaults to 'end'.

Pandas: Time series

Moving window functions

Función	Descripción
rolling_count	Returns number of non-NA observations in each trailing window.
rolling_sum	Moving window sum.
rolling_mean	Moving window mean.
rolling_media n	Moving window meadian.
rolling_std / rolling_var	Moving window variance / standard deviation. Uses (n-1) denominator.
rolling_skew / rolling_kurt	Moving window skewness (3rd moment) / kurtosis (4th moment).

Pandas: Time series

Moving window functions

Function	Description
rolling_min / rolling_max	Moving window minimum / maximum.
rolling_quantile	Moving window score at percentile / sample quantile.
rolling_corr / rolling_cov	Moving window correlation / covariance.
rolling_apply	Apply generic array function over a moving window.

Pandas: Time series

Moving window functions

Function	Description
ewma	Exponentially-weighted moving average.
ewmstd / ewmvar	Exponentially-weighted moving standard deviation / variance.
ewmcorr / ewmcov	Exponentially-weighted moving correlation / covariance.

Bibliography

- *Python for Data Analysis: Data Wrangling with Pandas, NumPy, and Python*. Wes McKinney. O'Reilly Media, 2017 (2^a edición).
- Python Data Science Handbook: Essential Tools for Working with Data. Jake VanderPlas. O'Reilly Media, 2016 (1^a edición).

High level tools

Contents

- Visualization: Seaborn.
- Machine learning: Scikit-Learn.
- Out-of-core computation: Dask.



Seaborn

- Seaborn is a graphics library that improves the design and appearance of Matplotlib plots and includes additional statistical methods:
 - It allows to use themes to unify plot aesthetics.
 - Adds functions to visualize and compare distributions on one and two variables.
 - Linear regression tools.
 - Functions for visualizing data matrices and clustering.
 - Plotting statistical time series.
- Similar to ggplot in R.



Scikit-Learn

- Extension to SciPy (**SciPy Toolkit**) focused on machine learning.
- Compatible with tabular data from other libraries:
 - NumPy
 - Pandas
 - SciPy.sparse (matrices dispersas).
- Includes algorithms for supervised / unsupervised learning: GLMs, SVMs, kNN, Bayes, decision trees, clustering, etc.
- Cross-validation.
- Grid search of optimal models.
- Parallelization.

Dask



- Limited but simple alternative to Spark.
- Restrictions to Big Data in Python:
 - Parallelization: restricted by the GIL.
 - Physical memory.
- Dask provides two main tools:
 - Dynamic task planifier (cluster management).
 - “Big data” collections: trivially perform out-of-core computing.
- Provides limited out-of-core support to other APIs:
 - Pandas
 - Scikit-Learn
- Allows to parallelize computations on a cluster.

Python: Other tools

- Statistical modeling:
 - Statsmodels.
- Big Data:
 - Hadoop.
 - Spark.
- Deep Learning:
 - TensorFlow.
 - PyTorch.
- GPU data science:
 - RAPIDS.

Bibliography

- Python Data Science Handbook: Essential Tools for Working with Data. Jake VanderPlas. O'Reilly Media, 2016 (1^a edición).
- Data Science with Python and Dask. Jesse C. Daniel. Manning Publications, 2019 (1^a edición).